

# Monitoring and Diagnostics

--- without OracleEnterpriseManager

Hemant K Chitale

<http://hemantoracledba.blogspot.com>

@HemantKChitale

SoS 14-July2015

## `whoami`

- Been a DBA / SME for more than 20years
- Began with V6. Worked on various Unix flavours, Linux, Windows. OPS and RAC
- Used OEM only for a couple of years in 2006-2007 (yes, the last time I used OEM was mid-2007)
- I will make these slides (with Notes) available via my blog and slideshare in a few days

I am not OEM savvy !

The notes sections in this presentation also provide links to some of my blog posts

## Privileges

- What privileges do you need ?
    - CREATE SESSION
    - SELECT\_CATALOG\_ROLE \*\*
    - EXECUTE ON DBMS\_XPLAN (present)
    - EXECUTE ON DBMS\_WORKLOAD\_REPOSITORY (optional)
  - What privileges do you NOT need ?
    - The DBA Role
    - SELECT ANY DICTIONARY
    - SELECT ANY TABLE
- \*\* Note : Direct Privileges required if using PLSQL

Only a small set of privileges are really required.

Also look at the OEM\_MONITOR role that has ANALYZE ANY , SELECT ANY, SELECT ANY DICTIONARY and ADVISOR privileges, none of which I use (although ADVISOR may be useful).

For a short note on the difference between SELECT ANY DICTIONARY and SELECT\_CATALOG\_ROLE see <http://hemantoracledba.blogspot.com/2014/02/the-difference-between-select-any.html>

SELECT ANY can be useful if you need to query data (e.g. To identify distribution patterns and skew) in a schema – but the owner might prefer to grant you SELECT privileges on only a subset of tables --- that sort of requirement comes in handy when doing Performance Tuning, which is outside of the scope of this presentation.

If you plan to use PLSQL (e.g. to schedule jobs to collect this monitoring information), you *will* need direct privileges on the underlying views. For example, on V\_\$SESSION, V\_\$SQL etc.

There are many System Privileges and Object level privileges that can be granted to Junior DBAs, Performance Analysts etc without having to grant the DBA role.

## Creating the User

```
SYSTEM>create user diag_mon identified by diag_mon;
User created.
SYSTEM>grant create session to diag_mon;
Grant succeeded.
SYSTEM>grant select_catalog_role to diag_mon;
Grant succeeded.
SYSTEM>select privilege, grantee from dba_tab_privs
  2 where table_name = 'DBMS_XPLAN';
PRIVILEGE                                GRANTEE
-----
EXECUTE                                PUBLIC
SYSTEM>connect sys/oracle as sysdba
Connected.
SYS>grant execute on dbms_workload_repository to diag_mon;
Grant succeeded.
SYS>
SYS>connect diag_mon/diag_mon
Connected.
DIAG_MON>
```

These are the privileges I need. I don't use ADVISORS but do use AWR.

## Some of the Views we'll look at here

- V\$SESSION
- V\$ACTIVE\_SESSION\_HISTORY (and DBA\_HIST\_ACTIVE\_SESS\_HISTORY)
- V\$SQL and V\$SQLSTATS
- V\$SESSION\_LONGOPS

See

```
select table_name, privilege from dba_tab_privs  
where grantee = 'SELECT_CATALOG_ROLE'
```

The fewer V\$ joins you need to make the better. Also, some V\$ views are preferable over others. Remember that Oracle does not provide the same read-consistency for V\$ (and X\$ !) views as for permanent tables and views. Joining V\$ views (to each other or to DBA\_% views) does not guarantee read consistency across the join.

## Active / Long Running Sessions

- V\$SESSION :
  - STATUS
  - LAST\_CALL\_ET
  - SEQ#, EVENT, STATE, SECONDS\_IN\_WAIT
  - Interpreting correlation between STATE, EVENT and SECONDS\_IN\_WAIT
  - LAST\_CALL\_ET and PLSQL and SQL
- V\$ACTIVE\_SESSION\_HISTORY

Columns that have been retrieved from V\$SESSION\_WAIT are now available in V\$SESSION. V\$SQL preferred over V\$SQLAREA because the latter does an aggregation (across all Child Cursors).

When I join V\$SESSION to V\$SQL, I join on both SQL\_ID and CHILD\_NUMBER.

Similarly, I prefer V\$SEGSTAT over V\$SEGMENT\_STATISTICS – the former is faster.

If you don't have the Diagnostic Pack for V\$ASH, you could sample V\$SESSION quickly with custom code.

# Active Running Session

```
DIAG_MON>1
  1 select username, sid, serial#, sql_id, status, last_call_et, seq#, event, state,
    seconds_in_wait
  2 from v$session
  3 where status = 'ACTIVE'
  4 and last_call_et > 4
  5* and type != 'BACKGROUND'
DIAG_MON>/

no rows selected

DIAG_MON>/

USERNAME          SID    SERIAL# SQL_ID          STATUS  LAST_CALL_ET    SEQ#
-----
EVENT              SECONDS_IN_WAIT                                STATE
-----
-----
HEMANT              9      3 gr4zw7dd73rlx ACTIVE          5      16400
direct path read                                WAITED SHORT TIME
0
```

What I intend to show is how to interpret the STATE. WAITED\_SHORT\_TIME means that it is *\*not\** currently in a Wait. The EVENT is the last wait, not current. So, although the session is ACTIVE, it is not in a Wait. It is most likely on CPU.

```

00:13:53 DIAG_MON>1
  1 select sql_id, status, last_call_et, seq#, event, state, seconds_in_wait
  2 from v$session
  3* where sid=195
00:14:35 DIAG_MON>/

SQL_ID          STATUS  LAST_CALL_ET    SEQ# EVENT                                STATE
-----
addfgac99vuxw ACTIVE          8      50735 db file scattered read        WAITED SHORT TIME
0

00:14:36 DIAG_MON>
00:14:46 DIAG_MON>/

SQL_ID          STATUS  LAST_CALL_ET    SEQ# EVENT                                STATE
-----
addfgac99vuxw ACTIVE          19     53225 free buffer waits                WAITING
0

00:14:47 DIAG_MON>
00:15:21 DIAG_MON>/

SQL_ID          STATUS  LAST_CALL_ET    SEQ# EVENT                                STATE
-----
addfgac99vuxw ACTIVE          53     60431 db file scattered read        WAITING
0

00:15:21 DIAG_MON>

```

Note the change to the WAITING on “free buffer waits”. That *is* the CURRENT Wait as at the time of the snapshot. (Similarly, the “db file scattered read” wait after that). So, at the bottom of this slide, the session’s current SQL has been active for 55 seconds and is currently in a multiblock read wait.



```
00:15:45 DIAG_MON>/
```

SQL_ID SECONDS_IN_WAIT	STATUS	LAST_CALL_ET	SEQ# EVENT	STATE
cdctxzddqb067 0	ACTIVE	78	518 db file scattered read	WAITING

```
00:15:46 DIAG_MON>
```

```
00:16:23 DIAG_MON>/
```

SQL_ID SECONDS_IN_WAIT	STATUS	LAST_CALL_ET	SEQ# EVENT	STATE
330q95smuwnv9 SHORT TIME	ACTIVE	116 0	11364 db file sequential read	WAITED

```
00:16:24 DIAG_MON>
```

Note how the \*current\* wait status can keep changing. Have you noted SEQ# incrementing ? That indicates that the Wait Events \*are\* changing.

## SQL called from PLSQL

- Interpreting SQL\_ID, STATUS, LAST\_CALL\_ET
- For example, for a DBMS\_STATS.GATHER\_TABLE\_STATS call, you will see the SQL\_ID keep changing from that of the DBMS\_STATS call to one of the child SQL calls that actually reads table / column / index data.
- The LAST\_CALL\_ET would be the time since the \*parent\* DBMS\_STATS call began !
- This also applies to PLSQL that calls SQL

Recursive SQLs called by DBMS\_STATS (or any PLSQL procedure) are at a depth level below. The top level is dep=0, the succeeding levels are 1 and beyond.

## SQL called from PSQL

```
HEMANT>exec dbms_stats.gather_table_stats('','LARGE_TABLE',method_opt=>'FOR ALL COLUMNS SIZE 250');
```

22:58:09 DIAG\_MON>-- running monitoring query  
22:58:10 DIAG\_MON>/

SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#	EVENT	STATE	SECONDS_IN_WAIT
135	65	98bht7550nwkr	ACTIVE	11	2652	db file scattered read	WAITING	0

22:58:14 DIAG\_MON>/

SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#	EVENT	STATE	SECONDS_IN_WAIT
135	65	52p6wm9z96c00	ACTIVE	21	7278	db file sequential read	WAITING	0

22:58:23 DIAG\_MON>/

SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#	EVENT	STATE	SECONDS_IN_WAIT
135	65	8y9qnq9m089b6	ACTIVE	25	20565	PX Deg: Execute Reply	WAITING	0

The GATHER\_TABLE\_STATS ran for 25seconds. The LAST\_CALL\_ET against SID=135 (HEMANT's session) was incremented across all the SQL calls, even though the SQL calls (at lower depths of 1 and below) were changing. Therefore, in this case, the LAST\_CALL\_ET is not for the SQL that was executing at that instant but for the calling PLSQL – the DBMS\_STATS.GATHER\_TABLE\_STATS call. So, when running PLSQL beware that LAST\_CALL\_ET may not reflect the current SQL !

(Note the SQL\_ID 8y9... statement – it was waiting on PQ slaves (I haven't shown those PQ slave sessions, but they do exist))

## SQL called from PLSQL

```

DIAG_MON>1
 1 select sql_id, sql_text from v$sql
 2* where sql_id in ('98bht7550nwkr','52p6wm9z96c00','8y9qng9m089b6','53nh88nc8mx1b') order by 1
DIAG_MON>/

SQL_ID
-----
SQL_TEXT
-----
-----
52p6wm9z96c00
SELECT topology FROM SDO_TOPO_METADATA_TABLE a, TABLE(a.Topo_Geometry_Layers) b WHERE b.owner =
:owner AND b.table_name = :tab

53nh88nc8mx1b
BEGIN dbms_stats.gather_table_stats('','LARGE_TABLE',method_opt=>'FOR ALL COLUMNS SIZE 250'); END;

8y9qng9m089b6
select /*+ parallel_index(t, "LARGE_TABLE_OWNER",4) dbms_stats cursor_sharing_exact
use_weak_name_resl dynamic_sampling(0) no_monitoring no_substr_pad no_expand
index_ffs(t,"LARGE_TABLE_OWNER") */ count(*) as nrw,count(distinct
sys_op_lbid(114855,'L',t.rowid))

98bht7550nwkr
/* SQL Analyze(0) */ select /*+ full(t) no_parallel(t) no_parallel_index(t) dbms_stats
cursor_sharing_exact use_weak_name_resl dynamic_sampling(0) no_monitoring no_substr_pad
*/to_char(count("OWNER")) ,to_char(substrb(dump(min("OWNER"),16,0,32),1,120)),to_cha

```

This slide is not necessarily part of this presentation. It is just to demonstrate that a PLSQL (the DBMS\_STATS.GATHER\_TABLE\_STATS in this case) can call multiple SQLs, at different recursive depth levels. (From dep=0 to dep=4 in this case). So, when you are monitoring V\$SQL\_ID in V\$SESSION, you might get a statement that is at a much lower depth.

# Where's the Catch ?

```
23:40:33 DIAG_MON>1
 1 select username, sid, serial#, sql_id, status, last_call_et, seq#, event, state, seconds_in_wait
 2 from v$session
 3 where username = 'HEMANT'
 4* and type != 'BACKGROUND'
23:40:35 DIAG_MON>/
```

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#	EVENT	STATE	SECONDS_IN_WAIT
HEMANT	9	3	gr4zw7dd73rlx	INACTIVE	492	17821	SQL*Net message from client	WAITING	492

```
23:40:35 DIAG_MON>
```

Note here that this session is WAITING on a message from client. Is it Idle ? Should the DBA ignore this ? Let's look at the next slide.

## Where's the Catch ?

```
23:40:35 DIAG_MON>
23:41:04 DIAG_MON>/
```

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#
EVENT						
HEMANT	9	3	9t7c0jy3v19pn	INACTIVE	0	24807
SQL*Net message from client					WAITING	0

```
23:41:07 DIAG_MON>/
```

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#
EVENT						
HEMANT	9	3	9t7c0jy3v19pn	INACTIVE	0	32663
SQL*Net message from client					WAITING	0

```
23:41:22 DIAG_MON>/
```

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#
EVENT						
HEMANT	9	3	9t7c0jy3v19pn	INACTIVE	0	64336
SQL*Net message from client					WAITING	0

Note SEQ# incrementing ? Even though the status is INACTIVE and the wait is "SQL\*Net message from client" (In theory an "idle wait") ?

What is really happening here is that the user or application server is running a query that is retrieving a large number of rows. Between every ARRAY FETCH, the server process has to wait for the client to acknowledge having received the rows – that is the "SQL\*Net message from client" wait.

Here the user says that his query is ACTIVE. And he is right ! The user is complaining that the query is long running but the actual Database Time the DBA sees is very low. Most of the time is being spent sending rows and waiting for the acknowledgement.

# Where's the Catch ?

23:42:21 DIAG\_MON>/

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#
-----						
EVENT					STATE	SECONDS_IN_WAIT
-----						
HEMANT	9	3	9t7c0jy3v19pn	INACTIVE	0	54223
SQL*Net message from client					WAITING	0

23:42:46 DIAG\_MON>/

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#
-----						
EVENT					STATE	SECONDS_IN_WAIT
-----						
HEMANT	9	3	9t7c0jy3v19pn	INACTIVE	4	34020
SQL*Net message from client					WAITING	4

...

Note the last wait at SEQ#34020. It is now 4 seconds ! Let's see the next slide.

## Where's the Catch ?

23:52:02 DIAG\_MON>/

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#
EVENT					STATE	SECONDS_IN_WAIT
HEMANT	9	3	9t7c0jy3v19pn	INACTIVE	560	34020
SQL*Net message from client					WAITING	560

23:52:03 DIAG\_MON>

23:52:26 DIAG\_MON>/

USERNAME	SID	SERIAL#	SQL_ID	STATUS	LAST_CALL_ET	SEQ#
EVENT					STATE	SECONDS_IN_WAIT
HEMANT	9	3	9t7c0jy3v19pn	INACTIVE	584	34020
SQL*Net message from client					WAITING	584

23:52:27 DIAG\_MON>

The SEQ# is still 34020. Now, this is truly an idle “SQL\*Net message from client” wait. The server is waiting for the client to send the next SQL command. It has finished sending the results of the last SQL statement. LAST\_CALL\_ET now reflect the time it is idle – because it is the time since the last call ended.



## Get Execution Plan (without SELECT ANY TABLE)

```
DIAG_MON>select * from table(dbms_xplan.display_cursor('gr4zw7dd73r1x'));
PLAN_TABLE_OUTPUT
```

```
-----
SQL_ID  gr4zw7dd73r1x, child number 0
-----
```

```
select count(*) from (select * from large_table minus select * from
another_large_table)
```

```
Plan hash value: 1789283771
-----
```

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT					126K(100)	
1	SORT AGGREGATE		1				
2	VIEW		4802K			126K (1)	00:25:14
3	MINUS						
4	SORT UNIQUE		4802K	448M	615M	126K (1)	00:25:14
5	TABLE ACCESS FULL	LARGE_TABLE	4802K	448M		18935 (1)	00:03:48
6	SORT UNIQUE		1	207		3 (34)	00:00:01
7	TABLE ACCESS FULL	ANOTHER_LARGE_TABLE	1	207		2 (0)	00:00:01

```
Note
```

```
-----
- dynamic sampling used for this statement (level=2)
```

```
24 rows selected.
```

```
DIAG_MON>
```

I don't have SELECT privilege on the underlying table in the HEMANT schema. Yet, I can get the execution plan. I don't need the SELECT privilege on the underlying table(s) or the SELECT ANY privilege or the DBA role to be able to do this.

# Check Query Tables

```
DIAG_MON>select owner, table_name, num_rows, last_analyzed
  2  from dba_tables
  3  where table_name in ('LARGE_TABLE','ANOTHER_LARGE_TABLE');
```

OWNER	TABLE_NAME	NUM_ROWS	LAST_ANAL
HEMANT	LARGE_TABLE	4802944	27-FEB-15
HEMANT	ANOTHER_LARGE_TABLE		

```
DIAG_MON>
```

Because I have SELECT\_CATALOG\_ROLE, I can query the underlying statistics.

# AWR

```
DIAG_MON>exec dbms_workload_repository.create_snapshot;
```

```
PL/SQL procedure successfully completed.
```

```
DIAG_MON>
```

```
DIAG_MON>@?/rdbms/admin/awrrpt {OR Use SQLDeveloper 4.0.1}
```

```
Extracted from the AWR Report :
```

Elapsed Time (s)	Elapsed Time Executions	per Exec (s)	%Total	%CPU	%IO	SQL Id
52.7	1	52.68	16.1	15.2	9.1	cx87ww8c8t206

Module: SQL\*Plus  
insert into another\_large\_table select \* from large\_table

I can create AWR snapshots because I have EXECUTE on DBMS\_WORKLOAD\_REPOSITORY. I don't need to be granted the DBA role.

I can use SQLDeveloper 4.0.1 to generate an AWR report without logging in to the server as "oracle"

# XPLAN from AWR

```
DIAG_MON>select * from table(dbms_xplan.display_awr('cx87mw8c8t206'));
```

```
PLAN_TABLE_OUTPUT
```

```
-----  
SQL_ID cx87mw8c8t206  
-----
```

```
insert into another_large_table select * from large_table
```

```
Plan hash value: 1101256009
```

```
-----  
| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |  
-----  
| 0 | INSERT STATEMENT | | | | 18935 (100) | |  
| 1 | LOAD TABLE CONVENTIONAL | | | | | |  
| 2 | TABLE ACCESS FULL | LARGE_TABLE | 4802K | 448M | 18935 (1) | 00:03:48 |  
-----
```

```
14 rows selected.
```

```
DIAG_MON>
```

I can retrieve historical execution plans captured by AWR.

## V\$SESSION\_LONGOPS

- CAVEAT ! It is based on \*operations\*, not on the SQL Execution Time. An SQL execution can actually consist of multiple operations.
- An SQL that consists of multiple steps in the execution plan can consist of multiple operations, each appearing independently in this view.
- Parallel Query operations are split by block ranges (or partitions) – so may appear repeatedly, once for each block range / partition
- Nested Loop Full Table Scans can appear repeatedly
- Entries in this view may not be cleared quickly, you may see entries for sessions / queries that have completed

Too many people on the Internet think that this view shows how long the query is running and how long it is expected to continue.

It shows the \*current operation\* not the whole SQL. An SQL Query can consist of multiple operations. Even Parallel Query can run different block ranges using multiple passes, each pass is a separate operation (and each PQ slave a separate session, so a separate row in this view).

## Session with Long Op ??

```
HEMANT>set time on
00:31:02 HEMANT>create table a_vl_table
00:31:18 2 as select * from large_table
00:31:38 3 union all
00:31:40 4 select * from large_table_2
00:31:46 5 union all
00:31:50 6 select * from large_table_3
00:31:55 7 /
```

Table created.

00:35:09 HEMANT>

PLAN\_TABLE\_OUTPUT

-----  
Plan hash value: xxxxxxxx

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	CREATE TABLE STATEMENT		14M	1346M	99106 (67)	00:19:50
1	LOAD AS SELECT	A_VL_TABLE				
2	UNION-ALL					
3	TABLE ACCESS FULL	LARGE_TABLE	4802K	448M	18935 (1)	00:03:48
4	TABLE ACCESS FULL	LARGE_TABLE_2	4802K	448M	18935 (1)	00:03:48
5	TABLE ACCESS FULL	LARGE_TABLE_3	4802K	448M	18935 (1)	00:03:48

12 rows selected.

HEMANT>

Here is an SQL that a user is executing. He is querying 3 different tables. Let's see in the next few slides if V\$SESSION\_LONGOPS shows the execution of the whole SQL or parts (operations) only.

```

00:32:38 DIAG_MON>1
  1 select sid, sql_plan_line_id, sql_plan_operation, opname, target, sofar,
  2 totalwork, units, to_char(start_time,'HH24:MI:SS') StartTime,
  3 elapsed_seconds, time_remaining, message
  4 from v$session_longops
  5 where sofar != totalwork
  6* order by start_time
00:32:38 DIAG_MON>/
no rows selected
00:32:39 DIAG_MON>/
no rows selected
00:32:41 DIAG_MON>/
      SID SQL_PLAN_LINE_ID SQL_PLAN_OPERATION
-----
OPNAME
-----
TARGET                                     SOFAR
-----
TOTALWORK UNITS          STARTTIM ELAPSED_SECONDS
-----
TIME_REMAINING
-----
MESSAGE
-----
          9          4 TABLE ACCESS
Table Scan
HEMANT.LARGE_TABLE_2                                25661
   69616 Blocks                                00:32:36          15
          26
Table Scan:  HEMANT.LARGE_TABLE_2: 25661 out of 69616 Blocks done

```

Look at the SQL\_PLAN\_LINE\_ID. This extract from V\$SESSION\_LONGOPS is for only one step in the Execution Plan. The SOFAR, ELAPSED\_SECONDS and TIME\_REMAINING are for that one operation – reading table HEMANT.LARGE\_TABLE\_2

The estimated time for this operation is 15+26 = 41seconds.

```

00:32:52 DIAG_MON>/

      SID SQL_PLAN_LINE_ID SQL_PLAN_OPERATION
-----
OPNAME
-----
TARGET                                     SOFAR
-----
TOTALWORK UNITS                          STARTTIM ELAPSED_SECONDS
-----
TIME_REMAINING
-----
MESSAGE
-----
      9          4 TABLE ACCESS
Table Scan
HEMANT.LARGE_TABLE_2                      36132
    69616 Blocks                          00:32:36          32
      30
Table Scan:  HEMANT.LARGE_TABLE_2: 36132 out of 69616 Blocks done

00:33:09 DIAG_MON>

```

The operation has been running for 32seconds and still needs another 30seconds (i.e. 62seconds in all, not the 42seconds estimated earlier). Oracle is continuously revising the estimated.



```

00:33:26 DIAG_MON>/

      SID SQL_PLAN_LINE_ID SQL_PLAN_OPERATION
-----
OPNAME
-----
TARGET                                     SOFAR
-----
TOTALWORK UNITS                          STARTTIM ELAPSED_SECONDS
-----
TIME_REMAINING
-----
MESSAGE
-----
      9              4 TABLE ACCESS
Table Scan
HEMANT.LARGE_TABLE_2                      64196
      69616 Blocks                          00:32:36          51
      4
Table Scan:  HEMANT.LARGE_TABLE_2: 64196 out of 69616 Blocks done

00:33:27 DIAG_MON>

00:33:33 DIAG_MON>/
no rows selected

00:33:36 DIAG_MON>

```

After the operation on Execution Plan Step 4 had completed, V\$SESSION\_LONGOPS stopped reporting this session. But is the session still active ?

```

00:33:43 DIAG_MON>/

      SID SQL_PLAN_LINE_ID SQL_PLAN_OPERATION
-----
OPNAME
-----
TARGET                                     SOFAR
-----
TOTALWORK UNITS                          STARTTIM ELAPSED_SECONDS
-----
TIME_REMAINING
-----
MESSAGE
-----
      9              5 TABLE ACCESS
Table Scan
HEMANT.LARGE_TABLE_3                      13235
  69616 Blocks                          00:33:31          12
      51
Table Scan:  HEMANT.LARGE_TABLE_3: 13235 out of 69616 Blocks done

00:33:44 DIAG_MON>

```

The session is still active. It is now on Execution Plan Step 5 – the next table in the SQL operation.

A database server process can do only 1 thing at a time. If it is querying LARGE\_TABLE\_2, it cannot also be querying LARGE\_TABLE\_3 at the same time. The retrieval of rows from LARGE\_TABLE\_3 is sequentially done later ! (Parallel Query is a way around the fundamental rule that a process can be doing only one thing at any time – PQ spawns multiple processes to do multiple things (reads from different block ranges and/or partitions of the same table) concurrently)

```
00:34:02 DIAG_MON>/

      SID SQL_PLAN_LINE_ID SQL_PLAN_OPERATION
-----
OPNAME
-----
TARGET                                     SOFAR
-----
TOTALWORK UNITS                          STARTTIM ELAPSED_SECONDS
-----
TIME_REMAINING
-----
MESSAGE
-----
      9              5 TABLE ACCESS
Table Scan
HEMANT.LARGE_TABLE_3                      40881
  69616 Blocks                          00:33:31      31
      22
Table Scan:  HEMANT.LARGE_TABLE_3: 40881 out of 69616 Blocks done

00:34:03 DIAG_MON>
```

In the previous slide, the estimate for the read from HEMANT.LARGE\_TABLE\_3 was (12 + 51) 63seconds. It is now 53seconds.

```

00:34:14 DIAG_MON>/

      SID SQL_PLAN_LINE_ID SQL_PLAN_OPERATION
-----
OPNAME
-----
TARGET                                     SOFAR
-----
TOTALWORK UNITS                          STARTTIM ELAPSED_SECONDS
-----
TIME_REMAINING
-----
MESSAGE
-----
      9              5 TABLE ACCESS
Table Scan
HEMANT.LARGE_TABLE_3                      48049
      69616 Blocks                      00:33:31      45
      20
Table Scan:  HEMANT.LARGE_TABLE_3: 48049 out of 69616 Blocks done
00:34:17 DIAG_MON>

```

The estimate has now changed to 65seconds.

For another example of misreading V\$SESSION\_LONGOPS on a DML that does a Full Table Scan see <http://hemantoracledba.blogspot.com/2009/01/when-not-to-use-vsessionlongops.html>

## Array Size

- Clients / ETL servers can use differing Array Sizes to fetch rows
- This means that the number of round-trips can vary
- The (assumed to be idle) wait event 'SQL\*Net message from client' can be an indicator

I had mentioned earlier that a query that sends multiple rows to a client / application server sends the rows in batches – based on the ARRAY Size. Search my blog for examples of ARRAYSIZE (and LINESIZE and PAGESIZE if using an SQLPlus Client) {I have a few different blogposts on this}

I have shown earlier how the “SQL\*Net message from client” isn’t always an Idle Event. The presence of this wait event, with increasing SEQ# can indicate array fetches.

## Array Size

```
select sql_id, sql_text, rows_processed, fetches
from v$sqlstats
where upper(sql_text) like
'SELECT * FROM HKC_TARGET_1 ORDER BY *'
order by 1
```

```
SQL_ID
```

```
-----
```

```
SQL_TEXT
```

```
-----
```

```
ROWS_PROCESSED    FETCHES
```

```
-----
```

```
0g29ksyksenxyw
```

```
select * from hkc_target_1 tt order by 2
```

```
1142966    11431
```

```
637hm0n25b5gh
```

```
select * from hkc_target_1 order by 1
```

```
1142966    1142967
```

Note the two queries retrieved the same number of rows. The elapsed time reported by the client would have included the “SQL\*Net message wait from client” wait event on the server for the multiple round trips. The FETCHES count indicates the round-trips. The first SQL used an ARRAYSIZE of 100, the second was doing a Row-By-Row FETCH. (The extra 1 FETCH is always present when you run an SQL, you’ll even see it in the trace file as a FETCH with 0 rows executed first).

## V\$ACTIVE\_SESSION\_HISTORY

- A snapshot is collected every second. Thus, NOT every Wait or ON\_CPU status is collected.
- To see if waits are increasing, also monitor SEQ#
- To track an SQL, look for the SQL\_ID appearing in consecutive snapshots for the same SESSION\_ID, SESSION\_SERIAL#
- Never try to add TIME\_WAITED from V\$A\_S\_H
- Note : Querying V\$A\_S\_H requires the Diagnostic Pack Licence. If you don't have access, you'll have to use V\$SESSION and sample it frequently
- V\$ASH availability is based on shared\_pool\_size and load, there is no guarantee how far back you can look

DBA\_HIST\_ACTIVE\_SESS\_HISTORY is a sample every 10seconds (not "1 in 10samples"). If I want to see the "distribution of a session or SQL over the CPU and wait events", I look at the number of samples, not a summation of TIME\_WAITED.

The composite key for an SQL execution within a session is SQL\_ID, SQL\_EXEC\_START, SQL\_EXEC\_ID

Remember : SQL Operations that completed between 2 snapshots (SAMPLE\_TIME) are *\*not\** captured !!

Here I present only a few examples of analysis using this view. There are many more useful columns like BLOCKING%, CURRENT%, QC% (e.g. I've seen people look at PGA\_ALLOCATED and TEMP\_SPACE\_ALLOCATED and have a query based on these as well).

Where a session has spent time (note : add filter by  
sample\_time)

```
DIAG_MON>1
 1 select session_state, event, count(*)
 2 from V$active_session_history
 3 where session_id=195
 4 - and sample_time between ' ' and ' '
 5 group by session_state, event
 6* order by 1,2
DIAG_MON>/
```

SESSION EVENT	COUNT (*)
ON CPU	2
WAITING cursor: pin S wait on X	3
WAITING db file sequential read	10
WAITING log file sync	1
WAITING read by other session	5

```
DIAG_MON>
```

I've not shown the filter by SAMPLE\_TIME here. It is a very short elapsed time of operations by Session 195. I can see that most of the samples indicate waiting on "db file sequential read" – more so than On CPU. I say "more samples" rather than "more time" as being more accurate.



### Distribution across SQLs (note : add filter by sample\_time if necessary)

```
DIAG_MON>1
 1 select sql_id, session_state, event, count(*)
 2 from V$active_session_history
 3 where session_id=195
 4 group by sql_id, session_state, event
 5 having count(*) > 2
 6* order by 1,4,2,3
```

DIAG\_MON>/

SQL_ID	SESSION EVENT	COUNT(*)
5ys3vrapmbx6w	WAITING direct path read	18
64mgk3gjr8pdk	ON CPU	8
83n9t3c9rsxfj	WAITING Data file init write	3
83n9t3c9rsxfj	WAITING buffer busy waits	10
83n9t3c9rsxfj	WAITING db file scattered read	23
83n9t3c9rsxfj	WAITING free buffer waits	26
83n9t3c9rsxfj	ON CPU	30
83n9t3c9rsxfj	WAITING log buffer space	76
calqgpqu9c2v4	WAITING read by other session	4
g0jvz8csyrtcf	WAITING db file sequential read	4

10 rows selected.

DIAG\_MON>

Here, the session ran multiple SQL statements. I can see the distribution of CPU and Wait Events amongst the different SQL.

By Module, rather than session (note : add filter by sample\_time)

```
DIAG_MON>1
 1 select sql_id, session_state, event, count(*)
 2 from V$active_session_history
 3 where module = 'SQL*Plus'
 4 group by sql_id, session_state, event
 5 having count(*) > 2
 6* order by 1,4,2,3
DIAG_MON>/
```

SQL_ID	SESSION EVENT	COUNT(*)
5ys3vrapmbx6w	WAITING direct path read	18
64mgk3gjr8pdk	ON CPU	8
83n9t3c9rsxfj	WAITING Data file init write	3
83n9t3c9rsxfj	WAITING buffer busy waits	10
83n9t3c9rsxfj	WAITING db file scattered read	23
83n9t3c9rsxfj	WAITING free buffer waits	26
83n9t3c9rsxfj	ON CPU	30
83n9t3c9rsxfj	WAITING log buffer space	76
	WAITING log file sync	12

9 rows selected.

DIAG\_MON>

Here, I filter for a MODULE, rather than a SESSION (I've not shown the filter by SAMPLE\_TIME). (Note : "log file sync" wait may not always show which the SQL\_ID was that was waiting on the Event).

## By SQL (which SQL is “busiest”)

```
DIAG_MON>1
 1 select sql_id, session_state, event, count(*)
 2 from V$active_session_history
 3 where
 4 sample_time > sysdate - 0.5/24
 5 and session_type = 'BACKGROUND'
 6 group by sql_id, session_state, event
 7 having count(*) > 2
 8* order by 1,4,2,3
DIAG_MON>/
```

SQL_ID	SESSION EVENT	COUNT(*)
5ys3vrapmbx6w	WAITING direct path read	18
64mgk3gjr8pdk	ON CPU	8
83n9t3c9rsxfj	WAITING Data file init write	3
83n9t3c9rsxfj	WAITING buffer busy waits	10
83n9t3c9rsxfj	WAITING db file scattered read	23
83n9t3c9rsxfj	WAITING free buffer waits	26
83n9t3c9rsxfj	ON CPU	30
83n9t3c9rsxfj	WAITING log buffer space	76
cvn54b7yz0s8u	WAITING db file sequential read	4
	WAITING log file sync	12

10 rows selected.

DIAG\_MON>

Here, I query for the last 30minutes.

## By SQL Plan Line

```
DIAG_MON>1
1 select sample_time, sql_id, sql_plan_line_id,session_state, event
2 from v$active_session_history
3 where session_id=140
4 and sample_time > sysdate-0.5/24
5* order by 1
```

SAMPLE_TIME	SQL_ID	SQL_PLAN_LINE_ID	SESSION	EVENT
---				
20-JUN-15 11.34.41.811 PM	8zk9mqxkr2jsq			ON CPU
20-JUN-15 11.37.30.274 PM	0f59859k2n07p	1	ON CPU	
20-JUN-15 11.39.20.646 PM	f90msdgm5858k	37	WAITING	direct path read
20-JUN-15 11.39.21.646 PM	f90msdgm5858k	2	ON CPU	
20-JUN-15 11.39.22.646 PM	f90msdgm5858k	37	WAITING	direct path read
...				
20-JUN-15 11.39.27.646 PM	f90msdgm5858k	37	ON CPU	
20-JUN-15 11.39.28.656 PM	f90msdgm5858k	37	WAITING	direct path read
...				
20-JUN-15 11.39.43.666 PM	f90msdgm5858k	37	WAITING	direct path read
20-JUN-15 11.41.20.212 PM	drna7u98myhhx	18	ON CPU	
20-JUN-15 11.41.21.212 PM	drna7u98myhhx	18	ON CPU	
20-JUN-15 11.41.22.212 PM	drna7u98myhhx	19	ON CPU	
20-JUN-15 11.41.23.212 PM	drna7u98myhhx	18	ON CPU	
...				
20-JUN-15 11.41.37.232 PM	drna7u98myhhx	19	ON CPU	

Here, I can also see, for each SQL, the time on each Step in the Execution Plan. So, I know which Step is likely to have accounted for more time (on the basis of the number of times it was sampled) !

## Distribution of Approx CPU usage – by Service

```
select s.name, count(*) On_CPU_Count,
       trunc((ratio_to_report(count(*)) over ())*100,0) percentage
from dba_hist_active_sess_history h, dba_services s
where h.service_hash=s.name_hash
and h.sample_time between to_date('&from_date','DD-MON-RR')
                      and to_date('&to_date','DD-MON-RR')
and h.session_state='ON CPU'
group by s.name
order by 3,1
/
```

NAME	ON_CPU_COUNT	PERCENTAGE
SYSS\$BACKGROUND	6351	27
SYSS\$USERS	16826	72

Distribution of CPU Usage (this is an approximation because it is based on a sample taken every 10seconds only !) This is based on number of occurrences in samples, not actual time spent on CPU. But we could approximate the one for the other.

## CPU and Waits for a User

```
select DECODE(On_CPU.Sample_Time,NULL,In_Wait.Sample_Time,On_CPU.Sample_Time) as Sample_Time, NVL(On_CPU.CNT,0) as On_CPU,
       NVL(In_Wait.CNT,0) as In_Wait
from
(
  select sample_time, count(*) CNT
  from
  (
    select h.sample_time, h.session_state, h.event
    from v$active_session_history h
    where 1=1
    and h.user_id='&&userid'
    and h.sample_time > sysdate-1/24
  )
  where session_state='ON CPU'
  group by sample_time
) On_CPU FULL OUTER JOIN
(
  select sample_time, count(*) CNT
  from
  (
    select h.sample_time, h.session_state, h.event
    from v$active_session_history h
    where 1=1
    and h.user_id='&&userid'
    and h.sample_time > sysdate-1/24
  )
  where session_state='WAITING'
  -- and event='db file sequential read'
  group by sample_time
) In_Wait
ON In_Wait.sample_time=On_CPU.sample_time
order by 1,2
/
```

To chart the number of sessions in ON\_CPU and in WAITING states. You could select for a specific wait event as well.

(SQLPLUS allows you to use SET COLSEP ',' and SET PAGESIZE 0 to spool to a CSV file)

Note : Information available in V\$ACTIVE\_SESSION\_HISTORY is limited by memory space allocated.

## Active Sessions Count

SAMPLE_TIME	ON_CPU	IN_WAIT
12-JUL-15 02.34.00.239 PM	1	0
12-JUL-15 02.34.15.329 PM	0	1
12-JUL-15 02.34.18.339 PM	3	0
12-JUL-15 02.34.19.339 PM	0	5
12-JUL-15 02.34.20.349 PM	5	0
12-JUL-15 02.34.21.359 PM	6	0
12-JUL-15 02.34.22.359 PM	7	1
12-JUL-15 02.34.23.369 PM	0	5
12-JUL-15 02.34.24.369 PM	0	5
12-JUL-15 02.34.25.399 PM	0	7
12-JUL-15 02.34.26.409 PM	0	7
12-JUL-15 02.34.27.409 PM	0	5
12-JUL-15 02.34.28.409 PM	0	4
12-JUL-15 02.34.29.419 PM	0	4
12-JUL-15 02.34.30.429 PM	0	5
12-JUL-15 02.34.31.429 PM	3	2
12-JUL-15 02.34.32.439 PM	0	5
12-JUL-15 02.34.33.439 PM	0	5
12-JUL-15 02.34.34.439 PM	0	5
12-JUL-15 02.34.35.449 PM	0	5
12-JUL-15 02.34.36.459 PM	0	5
12-JUL-15 02.34.37.459 PM	0	5

This is another way to represent Active Sessions.

## SQL Executions History

```
select  sq.snap_id snap_id,
        to_char(plan_hash_value) Plan_Hash_Value,
        executions_delta Executions,
        elapsed_time_delta/1000000 Total_Elapsed_Time_Seconds,
        elapsed_time_delta/1000000/decode(executions_delta,0,1,executions_delta)
          Elapsed_Time_Seconds_PerExec,
        disk_reads_delta/decode(executions_delta,0,1,executions_delta) Disk_Reads_PerExec,
        buffer_gets_delta/decode(executions_delta,0,1,executions_delta) Buffer_Gets_PerExec,
        rows_processed_delta/decode(executions_delta,0,1,executions_delta)
          Rows_Processed_PerExec
from    dba_hist_sqlstat sq, dba_hist_snapshot ss
where
sq.dbid=ss.dbid
and sq.snap_id=ss.snap_id
and
sql_id = '&sql_id'
--and executions_delta <> 0
order by snap_id
/
```

**Note :** Requires Diagnostic Pack Licence

Extract previous occurrences of an SQL from AWR history



## SQL Executions History

Snap_ID	Plan_Hash_Valu	EXECUTIONS	TOTAL_ELAPSED_TIME_SECS	ELAPSED_TIME_SECS_PEREXEC
1087	1586852085	3	65	
22	144,175	187,402	768,094	
1088	1586852085	1	98	
98	782,464	930,067	3,168,456	
1089	1586852085	2	72	
36	198,984	330,982	676,449	

1087	1586852085	3	65
22	144,175	187,402	768,094
1088	1586852085	1	98
98	782,464	930,067	3,168,456
1089	1586852085	2	72
36	198,984	330,982	676,449

The same SQL has had varying numbers of Buffer Gets and Rows Processed. For example, in Snapshot 1088, there was 1 execution for 3million rows. The other two snapshots had 2 or 3 executions for 676K or 768K rows each.

Note : PLAN\_HASH\_VALUE does not change if ROWS or COST changes. See <http://hemantoracledba.blogspot.com/2014/03/plan-hashvalue-remains-same-for-same.html> (So, two queries with the same P\_H\_V don't have to have the same expected runtime !)

## CPU Consumption by Resource Consumer Group

```
DIAG_MON>1
 1 select to_char(begin_time,'DD-MON HH24:MI') From_time, to_char(end_time,'DD-MON HH24:MI')
    To_time,
 2 consumer_group_name,
 3 60 * (select value from v$parameter where name = 'cpu_count') max_db_sec,
 4 cpu_consumed_time/1000 consumed_sec
 5 from v$srcmgtmetric_history
 6* order by begin_time, consumer_group_name
DIAG_MON>/
```

**Note :** View shows statistics only for the past 1 hour

FROM_TIME	TO_TIME	CONSUMER_GROUP_NAME	MAX_DB_SEC	CONSUMED_SEC
11-JUL 23:56	11-JUL 23:57	OTHER_GROUPS	240	.589
11-JUL 23:56	11-JUL 23:57	FX_QUERY_USER	240	12.55
11-JUL 23:56	11-JUL 23:57	_ORACLE_BACKGROUND_GROUP_	240	0
11-JUL 23:57	11-JUL 23:58	OTHER_GROUPS	240	.02
11-JUL 23:57	11-JUL 23:58	FX_QUERY_USER	240	10.085
11-JUL 23:57	11-JUL 23:58	_ORACLE_BACKGROUND_GROUP_	240	0
11-JUL 23:58	11-JUL 23:59	OTHER_GROUPS	240	.016
11-JUL 23:58	11-JUL 23:59	FX_QUERY_USER	240	18.993
11-JUL 23:58	11-JUL 23:59	_ORACLE_BACKGROUND_GROUP_	240	0
11-JUL 23:59	12-JUL 00:00	OTHER_GROUPS	240	1.527
11-JUL 23:59	12-JUL 00:00	FX_QUERY_USER	240	0
11-JUL 23:59	12-JUL 00:00	_ORACLE_BACKGROUND_GROUP_	240	0

If Resource Manager is implemented, I can also look at the CPU usage by each Consumer Group.

**Note :** When “Virtual CPUs” (as in VMs or HyperThreading) are used, the CPU Count may be larger than the number of actual cores so CPU time may be misleading relative to the actual number of cores.

## Parallel Query Usage

```

DIAG_MON>1
 1 select to_char(sysdate,'DD-MON HH24:MI'), p.qcsid, s.sql_id, p.req_degree, p.degree, count(*)
 2 from v$px_session p, v$session s
 3 where p.sid=s.sid
 4 and p.serial#=s.serial#
 5 and p.req_degree is not null
 6* group by to_char(sysdate,'DD-MON HH24:MI'), p.qcsid, s.sql_id, p.req_degree, p.degree
DIAG_MON>
DIAG_MON>/
TO_CHAR(SYSDATE,'DD-M   QCSID SQL_ID          REQ_DEGREE    DEGREE    COUNT(*)
-----
12-JUL 00:15           135 20mwzfa97gczs          4          4          8

DIAG_MON>/
TO_CHAR(SYSDATE,'DD-M   QCSID SQL_ID          REQ_DEGREE    DEGREE    COUNT(*)
-----
12-JUL 00:17           18 4k9ruvz7nmuhj          4          4          8
12-JUL 00:17          135 2f8xqbm44zxnn          4          4          4

DIAG_MON>/
TO_CHAR(SYSDATE,'DD-M   QCSID SQL_ID          REQ_DEGREE    DEGREE    COUNT(*)
-----
12-JUL 00:17           135 20mwzfa97gczs          4          4          8

DIAG_MON>/
TO_CHAR(SYSDATE,'DD-M   QCSID SQL_ID          REQ_DEGREE    DEGREE    COUNT(*)
-----
12-JUL 00:17           135 dfpjydp33tvnr          4          4          4

```

I can see how many PQs are running (look at QCSID as the Query Co-ordinator). I can see if each QC did get the actual degree (DEGREE) that was requested (REQ\_DEGREE). So, if the server doesn't have enough PARALLEL\_MAX\_SERVERS, the actual degree may be less than the requested degree.

Thus, at 00:15, there was 1 query requesting DoP=4, getting Dop=4 but actually running with 8 PQ slaves (Parallel Execution Servers). Why, because a query may take 2 (or, in very rare cases, more) Slave Sets.

At 00:17, there were two queries from two different sessions executing with a total of 12 PX servers.

See my blog posts on PX Servers.

# Transactions

- Look at START\_TIME, USED\_UBLK and USED\_UREC in V\$TRANSACTION. These do NOT map to Table Rows (or Rows + Index Entries)
- A Transaction can consist of multiple SQL calls
- You cannot identify which SQL in the transaction generated how much Undo
- USED\_UREC, USED\_UBLK start reducing when a rollback begins (user gets message only when rollback is completed)
- (join V\$TRANSACTION to V\$SESSSION on T.ADDR=S.TADDR)
- Direct Path operations (INSERT APPEND or PARALLEL) generate only 1 Undo Record

Very important : UNDO RECORDS is NOT the same (or same size) as Table Rows or Index Entries.

## Active Transaction

```
00:28:14 DIAG_MON>1
 1  select s.sid, s.serial#, s.username, substr(s.program,1,10) Prognm,
 2  t.xidusn, t.used_ublk, t.used_urec,t.start_time,
 3  s.last_call_et, s.sql_id Current_SQL_ID_if_available
 4  from v$session s, v$transaction t
 5* where s.taddr=t.addr
00:28:14 DIAG_MON>/
```

SID	SERIAL#	USERNAME	PROGNM	XIDUSN	USED_UBLK	USED_UREC	START_TIME
LAST_CALL_ET	CURRENT_SQL_I						
135	173	HEMANT	sqlplus@lo	9	5361	184675	07/12/15 00:27:00
74	gaplj5qttgc62						

```
00:28:15 DIAG_MON>
0:28:48 DIAG_MON>/
```

SID	SERIAL#	USERNAME	PROGNM	XIDUSN	USED_UBLK	USED_UREC	START_TIME
LAST_CALL_ET	CURRENT_SQL_I						
135	173	HEMANT	sqlplus@lo	9	5069	174809	07/12/15 00:27:00
196							

```
00:30:18 DIAG_MON>00:30:18 DIAG_MON>
```

You can check the distribution of transactions across different Undo Segments.

Remember : The SQL\_ID is only the current SQL. A transaction can consist of multiple DMLs, including SELECT queries ! The current SQL may be a SELECT but the transaction may have 1 or 100 or 1000 INSERT/UPDATE/DELETE statements before this.

The transaction started rolling back after 00:28:15. (Also note that Current SQL\_ID is not always available)

## Instance and Session Statistics and Events

- Use V\$SYSSTAT and V\$SESSTAT (joined to V\$STATNAME) for Cumulative Instance / Session Statistics.
- Use V\$SYSTEM\_EVENT and V\$SESSION\_EVENT for Cumulative Instance / Session Events (Waits)
- It is possible to do a query that Union of the STATS and EVENT views. (“CPU used by this session” is in the STATS view)
- I use this for a “Profile” view somewhat similar to the header information in AWR (\*without\* the “Hit Ratios”)

I can compare CPU time with Wait Events time. I can identify the top Wait. I can filter for statistics or waits that I am particularly interested in.

## Profiling Waits and Statistics

```
DIAG_MON>1
1  select 'Event : ' || event Event_or_Stat, time_waited, total_waits,
   total_timeouts, average_wait
2  from v$system_event
3  where
4      event not like 'rdbms ipc message%'
5  and event not like 'SQL*Net%'
6  and event not like 'pipe get%'
7  and event not like 'pmon timer%'
8  and event not like 'smon timer%'
9  union
10 select 'Statistic : ' || name a, value b, null c, null d, null e
11 from v$sysstat
12 where name in ('CPU used by this session')
13 union
14 select 'Statistic : ' || name a, null b, value c, null d, null e
15 from v$sysstat
16 where name in
17 ('consistent changes','consistent gets','db block gets', 'db block
   changes','physical reads',
18 'redo blocks written', 'redo entries', 'redo size', 'user calls', 'user commits')
19* order by 1
DIAG_MON>
```

Event Wait Time and CPU used are in CentiSeconds

I can select which Statistics and which Wait Events are of interest to me for this particular database instance. (Different application profiles may have different interesting Statistics and Waits !)

This gives me a “Profile” view of the Instance

## Instance Waits and Stastics

EVENT_OR_STAT	TIME_WAITED	TOTAL_WAITS	TOTAL_TIMEOUTS	AVERAGE_WAIT
Event : ARCH wait for process	202	2	2	101.16
Event : db file parallel read	99	41	0	2.41
Event : db file scattered read	4282	61863	0	.07
Event : db file sequential rea	63417	203607	0	.31
Event : db file single write	1	32	0	.03
Event : direct path read	159831	936937	0	.17
Event : direct path read temp	2	496	0	0
Event : log buffer space	571	20	0	28.57
Statistic : CPU used by this s	63995			
Statistic : consistent changes		17863		
Statistic : consistent gets		16648379		
Statistic : db block changes		920998		
Statistic : db block gets		762072		
Statistic : physical reads		15934919		
Statistic : redo blocks writte		312956		
Statistic : redo entries		456759		
Statistic : redo size		154028216		
Statistic : user calls		226709		
Statistic : user commits		1074		
sum		24803993		

Note : Time\_Waited and CPU used are both in Centi-seconds.

Here I can compare CPU time with time on major wait events. I can also look at major statistics that I can filter.