

Oracle Diagnostics

Hemant K Chitale

Hemant K Chitale

- whoami ?
- Oracle 5 to Oracle 10gR2 : DOS, Xenix, 8 flavours of Unix, Linux, Windows
- Financial Services, Govt/Not-for-Profit, ERP, Custom
- Production Support, Consulting, Development
- A DBA, not a Developer
- Product Specialist, Standard Chartered Bank
- [My Oracle Blog](http://hemantoracledba.blogspot.com) *http://hemantoracledba.blogspot.com*

Locks and Lock Trees

- Row Locks are Enqueues
- They serialise access to rows
- A transaction may hold Row Locks on multiple rows – this is represented as a single entry in V\$TRANSACTION but single or multiple entries in the ITL slots in various table / index blocks
- ITLs allow different transactions to lock different rows in the same block concurrently.
- Lock Trees are multiple sessions waiting “in order”, with potentially more than one session waiting on the same row lock

A Lock Tree :

Script “utllockt.sql” (in \$ORACLE_HOME/rdbms/admin) can provide a tree-like diagram.

```
* This script prints the sessions in the system that are waiting for
* locks, and the locks that they are waiting for. The printout is tree
* structured. If a sessionid is printed immediately below and to the right
* of another session, then it is waiting for that session. The session ids
* printed at the left hand side of the page are the ones that everyone is
* waiting for.
*
* For example, in the following printout session 9 is waiting for
* session 8, 7 is waiting for 9, and 10 is waiting for 9.
*
* WAITING_SESSION  TYPE MODE REQUESTED      MODE HELD          LOCK ID1 LOCK ID2
* -----
* 8                NONE None                None                0          0
* 9                TX  Share (S)          Exclusive (X)       65547      16
* 7                RW  Exclusive (X)      S/Row-X (SSX)      33554440   2
* 10               RW  Exclusive (X)      S/Row-X (SSX)      33554440   2
*
* The lock information to the right of the session id describes the lock
* that the session is waiting for (not the lock it is holding).
```

The script can be enhanced to provide more session information. The script uses DDLs to drop and create temp tables – so another enhancement would be to have those tables created in advance as GTTs and only populated and queried by the script

Here is a query to list Lock Holders and Waiters:

```
select s.blocking_session, to_number(s.sid) Waiting_Session, s.event, s.seconds_in_wait,
p.pid,
    p.spid "ServerPID", s.process "ClientPID",
    s.username, s.program, s.machine, s.osuser, s.sql_id,
    substr(sq.sql_text,1,75) SQL
from v$sql sq, v$session s, v$process p
where s.event like 'enq: TX%'
and s.paddr=p.addr
and s.sql_address=sq.address
and s.sql_hash_value=sq.hash_value
and s.sql_id=sq.sql_id
and s.sql_child_number=sq.child_number
union all
select s.blocking_session, to_number(s.sid) Waiting_Session, s.event, s.seconds_in_wait,
p.pid,
    p.spid "ServerPID", s.process "ClientPID",
    s.username, s.program, s.machine, s.osuser, s.sql_id,
    substr(sq.sql_text,1,75) SQL
from v$sql sq, v$session s, v$process p
where s.sid in (select distinct blocking_session from v$session where event like 'enq:
TX%')
and s.paddr=p.addr
and s.sql_address=sq.address(+)
and s.sql_hash_value=sq.hash_value(+)
and s.sql_id=sq.sql_id(+)
and s.sql_child_number=sq.child_number(+)
order by 1 nulls first, 2
/
```

This method does NOT require any temporary tables !

Example 1 :

Two separate sessions attempting to update the same row :

```
SQL> connect ABC_DBA/ABC_DBA_123
```

```
Connected.
```

```
SQL> update hemant.test_row_lock set content = 'Another' where  
pk=1;
```

```
1 row updated.
```

```
SQL>
```

```
SQL> connect hemant/hemant
```

```
Connected.
```

```
SQL> update test_row_lock set content = 'First' where pk=1;
```

..... now waiting

BLOCKING_SESSION WAITING_SESSION EVENT
SECONDS_IN_WAIT PID

ServerPID ClientPID USERNAME PROGRAM

MACHINE OSUSER
SQL_ID

SQL

17 SQL*Net message from client

82 27

13788 13770 ABC_DBA

sqlplus@localhost.localdomain (TNS V1-V3)

localhost.localdomain oracle

17

26 enq: TX - row lock contention

52 19

13791 3449 HEMANT

sqlplus@localhost.localdomain (TNS V1-V3)

localhost.localdomain oracle

fuwn3bnuh2axg

update test_row_lock set content = 'First' where pk=1

SQL>

Findings for Ex. 1 :

Session 17 (ABC_DBA) has no 'BLOCKING_SESSION'. It isn't waiting. It is the Blocker itself.

Session 26 (HEMANT) as a 'WAITING_SESSION' presents Session 17 as being the 'BLOCKING_SESSION'.

Session 26 is waiting on a Row Lock while currently running "update
test_row_lock set content = 'First' where pk=1"

Example 2 :

Table Definition :

```
SQL> create table test_unique_insert_row_lock  
(col_1 number , col_2 varchar2(5));
```

Table created.

```
SQL> create unique index t_u_i_r_l_unq_ndx on  
test_unique_insert_row_lock (col_1);
```

Index created.

```
SQL>
```

Example 2 (contd):

```
SQL> connect ABC_DBA/ABC_DBA_123
```

```
Connected.
```

```
SQL> insert into hemant.test_unique_insert_row_lock values  
(1, 'Sn.1');
```

```
1 row created.
```

```
SQL>
```

```
SQL> connect hemant/hemant
```

```
Connected.
```

```
SQL> insert into hemant.test_unique_insert_row_lock values  
(1, 'Sn.2');
```

```
..... now waiting .....
```

BLOCKING_SESSION WAITING_SESSION EVENT
SECONDS_IN_WAIT PID

ServerPID
PROGRAM

ClientPID

USERNAME

MACHINE
SQL_ID

OSUSER

SQL

1 SQL*Net message from client

115 22

13007

12813

ABC_DBA

sqlplus@localhost.localdomain (TNS V1-V3)

localhost.localdomain

oracle

1

36 enq: TX - row lock contention

80 25

13009

12971

HEMANT

sqlplus@localhost.localdomain (TNS V1-V3)

localhost.localdomain

oracle

166dv336yhxua

insert into hemant.test_unique_insert_row_lock values (1,'Sn.2')

2 rows selected.

SQL>

Findings for Ex 2 :

Session 1 (ABC_DBA) has no blocker

Session 36 (HEMANT) is waiting on Session 1

So, even a Duplicate Value on INSERT actually “waits” on the INSERT that has not committed. The row inserted by ABC_DBA is not visible to HEMANT (because ABC_DBA has not yet committed) but HEMANT’s INSERT waits on the Row Lock !