

Oracle Diagnostics

Hemant K Chitale

Hemant K Chitale

- whoami ?
- Oracle 5 to Oracle 10gR2 : DOS, Xenix, 8 flavours of Unix, Linux, Windows
- Financial Services, Govt/Not-for-Profit, ERP, Custom
- Production Support, Consulting, Development
- A DBA, not a Developer
- [My Oracle Blog](http://hemantoracledba.blogspot.com) *http://hemantoracledba.blogspot.com*

Explain Plans -- simple

- Explain Plan is a method of displaying the *expected* OR *actual* SQL Execution Plan
- Since 9i, Oracle has provided the DBMS_XPLAN package with various procedures

Method 1 : Without Executing the Query

- NOT to be used if the query has Binds – particularly because “Explain is blind to Binds”
- Use “EXPLAIN PLAN FOR Query” followed by “DBMS_XPLAN.DISPLAY”.

Given this query :

```
select sale_id, cust_id, remarks
from sales where
sale_date between to_date('01-NOV-10','DD-MON-RR')
                and to_date('04-NOV-10','DD-MON-RR')
```

```
SQL> explain plan for
  2  select sale_id, cust_id, remarks
  3  from sales where
  4  sale_date between to_date('01-NOV-10','DD-MON-RR')
  5                      and to_date('04-NOV-10','DD-MON-RR')
  6  /
```

Explained.

```
SQL> select * from table(dbms_xplan.display);
PLAN_TABLE_OUTPUT
```

```
-----
-----
Plan hash value: 1231079358
```

```
-----
| Id  | Operation                | Name   | Rows  | Bytes | Cost (%CPU)| Time     |
-----
|  0  | SELECT STATEMENT          |        |  3000 |  166K |    244   (2)| 00:00:03 |
|*  1  |   FILTER                  |        |       |       |           |          |
|*  2  |    TABLE ACCESS FULL    | SALES  |  3000 |  166K |    244   (2)| 00:00:03 |
-----
```

```
Predicate Information (identified by operation id):
```

```
-----
  1 - filter(TO_DATE('01-NOV-10','DD-MON-RR')<=TO_DATE('04-NOV-10','DD-
        MON-RR'))
  2 - filter("SALE_DATE"<=TO_DATE('04-NOV-10','DD-MON-RR') AND
        "SALE_DATE">=TO_DATE('01-NOV-10','DD-MON-RR'))
```

17 rows selected.

Understand the components :

Plan hash value: 1231079358

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		3000	166K	244 (2)	00:00:03
* 1	FILTER					
* 2	TABLE ACCESS FULL	SALES	3000	166K	244 (2)	00:00:03

Predicate Information (identified by operation id):

- 1 - filter(TO_DATE('01-NOV-10','DD-MON-RR')<=TO_DATE('04-NOV-10','DD-MON-RR'))
- 2 - filter("SALE_DATE"<=TO_DATE('04-NOV-10','DD-MON-RR') AND "SALE_DATE">=TO_DATE('01-NOV-10','DD-MON-RR'))

Every Execution Plan has a Hash Value. (Just as every SQL has a Hash Value and SQL_ID). We'll see later where the Hash Value is an important clue.

Step 2 is indented --- we normally think that it is executed before Step 1.

A Filter for SALE_DATE between two dates is applied when doing a FullTableScan ("TABLE ACCESS FULL" and "filter") at Step 2. Oracle expects to return 3000 rows after applying the "filter" to the FullTableScan. (The Explain Plan shows the number of rows expected to be returned by the step, not the number of rows that the FTS will read {which is actually 100,000 !}). These 3000 rows will be 166KBytes to be returned to the client (SQLPlus session).

Step 1 is a filter that Oracle applies for validation.

What is the filter in Step 1 ? Why does Oracle “apply it for validation” ?

```
SQL> 1
      2 explain plan for
      3 select sale_id, cust_id, remarks
      4 from sales where
      5 sale_date between to_date('01-NOV-10','DD-MON-RR')
      6 and to_date('25-OCT-10','DD-MON-RR')
SQL> /
```

Explained.

```
SQL> select * from table(dbms_xplan.display);
PLAN_TABLE_OUTPUT
```

```
-----
-----
Plan hash value: 1231079358
-----
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	57	244 (2)	00:00:03
* 1	FILTER					
* 2	TABLE ACCESS FULL	SALES	1	57	244 (2)	00:00:03

```
-----
Predicate Information (identified by operation id):
-----
```

- 1 - filter(TO_DATE('01-NOV-10','DD-MON-RR')<=TO_DATE('25-OCT-10','DD-MON-RR'))
- 2 - filter("SALE_DATE"<=TO_DATE('25-OCT-10','DD-MON-RR') AND "SALE_DATE">=TO_DATE('01-NOV-10','DD-MON-RR'))

```
SQL>
```



```
SQL> set autotrace on
SQL> select sale_id, cust_id, remarks
       2  from sales where
       3  sale_date between to_date('01-NOV-10','DD-MON-RR')
       4                        and to_date('25-OCT-10','DD-MON-RR')
       5  /
```

no rows selected

Execution Plan

Plan hash value: 1231079358

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	57	244 (2)	00:00:03
* 1	FILTER					
* 2	TABLE ACCESS FULL	SALES	1	57	244 (2)	00:00:03

Predicate Information (identified by operation id):

- ```
1 - filter(TO_DATE('01-NOV-10','DD-MON-RR')<=TO_DATE('25-OCT-10','DD-
 MON-RR'))
2 - filter("SALE_DATE"<=TO_DATE('25-OCT-10','DD-MON-RR') AND
 "SALE_DATE">=TO_DATE('01-NOV-10','DD-MON-RR'))
```

#### Statistics

```
1 recursive calls
0 db block gets
0 consistent gets
0 physical reads
0 redo size
409 bytes sent via SQL*Net to client
408 bytes received via SQL*Net from client
1 SQL*Net roundtrips to/from client
```

SQL>

You can now see that the FILTER in step 1 was a method for validation.

Since

```
filter(TO_DATE('01-NOV-10','DD-MON-RR')<=TO_DATE('25-OCT-10','DD-MON-RR'))
```

will be **\*FALSE\*** (because '01-NOV-10' is **\*not less than\*** '25-OCT-10'), the next step (Step 2) does not get executed at all.

That is why the AUTOTRACE shows 0 block gets

Statistics

-----

```
1 recursive calls
0 db block gets
0 consistent gets
0 physical reads
```

inspite of Step 2 supposedly being “TABLE ACCESS FULL” !

Therefore, when you see a FILTER as a step in an Execution Plan, remember that it may be a “logic filter”. Oracle may be using it to decide if the “indented” step below it needs to be executed at all.

If this FILTER returns FALSE, the “indented” “child” step is NOT executed.

## Method 2 : Without Query Execution

- This is useful to “validate” an Execution Plan
- This is an alternative to executing with SQL Tracing enabled and writing to a trace file
- This method uses the Hint “GATHER\_PLAN\_STATISTICS”. (Another option is to set the session parameter STATISTICS\_LEVEL to “ALL” just before execution of the SQL).

## Actually executing the query, with the Hint added :

```
SQL> select /*+ gather_plan_statistics */ sale_id, cust_id, remarks
 2 from sales where
 3 sale_date between to_date('01-NOV-10','DD-MON-RR')
 4 and to_date('04-NOV-10','DD-MON-RR')
 5 /
```

| SALE_ID | CUST_ID | REMARKS                                  |
|---------|---------|------------------------------------------|
| 24099   | 40      | BIK9SBLPJKGKA8UINWX20064KAD210CMW4Z7AZUL |
| 24100   | 44      | TXE1BTQM1631FJVTYCUBYBGOFRL7O32QVG20ZV6C |
| .....   | .       |                                          |
| .....   | .       |                                          |
| 27097   | 77      | 11KBENL0XE4T2OXAWXF9DAW2HHSQBG696BHJ5F79 |
| 27098   | 17      | MHEK21ILW79EHUIJC1Q3RA4PLC844K2KXNXCYL3E |

3000 rows selected.

Elapsed: 00:00:04.20

SQL>

```
SQL> select * from table(dbms_xplan.display_cursor('','','ALLSTATS LAST'));
```

```
PLAN_TABLE_OUTPUT
```

```

SQL_ID 2xk5b0ypks53n, child number 0

```

```
select /*+ gather_plan_statistics */ sale_id, cust_id, remarks from
sales where sale_date between to_date('01-NOV-10','DD-MON-RR')
and to_date('04-NOV-10','DD-MON-RR')
```

```
Plan hash value: 1231079358
```

```

| Id | Operation | Name | Starts | E-Rows | A-Rows | A-Time | Buffers | Reads |

0	SELECT STATEMENT		1		3000	00:00:00.24	1064	835
* 1	FILTER		1		3000	00:00:00.24	1064	835
* 2	TABLE ACCESS FULL	SALES	1	3000	3000	00:00:00.15	1064	835

```

```
Predicate Information (identified by operation id):

```

```
1 - filter(TO_DATE('01-NOV-10','DD-MON-RR')<=TO_DATE('04-NOV-10','DD-MON-RR'))
2 - filter(("SALE_DATE"<=TO_DATE('04-NOV-10','DD-MON-RR') AND
"SALE_DATE">=TO_DATE('01-NOV-10','DD-MON-RR')))
```

I use the DISPLAY\_CURSOR procedure.

The first two Parameters are SQL\_ID and CHILD\_NO. The " (NULL) values passed tell Oracle to evaluate for the SQL just (previously) executed in my \*current\* session.

The last parameter is for display format. 'ALLSTATS LAST' is to show all statistics for the Last execution of the same SQL.

| Id  | Operation         | Name  | Starts | E-Rows | A-Rows | A-Time      | Buffers | Reads |
|-----|-------------------|-------|--------|--------|--------|-------------|---------|-------|
| 0   | SELECT STATEMENT  |       | 1      |        | 3000   | 00:00:00.24 | 1064    | 835   |
| * 1 | FILTER            |       | 1      |        | 3000   | 00:00:00.24 | 1064    | 835   |
| * 2 | TABLE ACCESS FULL | SALES | 1      | 3000   | 3000   | 00:00:00.15 | 1064    | 835   |

Predicate Information (identified by operation id):

```

1 - filter(TO_DATE('01-NOV-10','DD-MON-RR')<=TO_DATE('04-NOV-10','DD-MON-RR'))
2 - filter(("SALE_DATE"<=TO_DATE('04-NOV-10','DD-MON-RR') AND
 "SALE_DATE">=TO_DATE('01-NOV-10','DD-MON-RR'))

```

Here Oracle shows the Expected-Rows from each Step versus the Actual-Rows. (We already know the Step 1 for this plan is a “validation” step so we ignore it). Step 2 was expected to return 3000 rows and did return 3000 rows (after applying the filters on all the 100,000 rows returned by the FullTableScan). The Actual-Time and Buffer Gets and Physical Reads (number of Oracle Blocks, not number of Read Calls to the OS) are also displayed.

In a complex Execution Plan, the A-Rows, A-Time and Buffers (and, possibly, Reads) columns are the ones to pay attention to – they will indicate whether the expected cardinality for the particular step was correct (if A-Rows is different from E-Rows) and how long the step took.

NOTE : When you have a query that is a join of two tables using a Nested Loop, you might mis-read the E-Rows and A-Rows. You will have to pay attention to another column “Starts”.

## Method 3 : From AWR Historical Information

- This is useful to monitor a query's execution profile over time
- It can be done only if the SQL has been captured by AWR
- Remember : AWR only captures the Top 'N' SQLs present in the Shared Pool when a Snapshot is created. If the SQL had been invalidated or aged out or was not in the Top 'N', at the time of the Snapshot, it wouldn't have been captured by AWR !

NOTE : This output spans this slide \*and\* the next two slides :

```
SQL> select * from table(dbms_xplan.display_awr('3pat1z2qx4gyg')) ;
```

```
PLAN_TABLE_OUTPUT
```

```



```



SQL\_ID 3pat1z2qx4gyg

```
select sale_id, cust_id, remarks from sales where sale_date between
to_date('01-NOV-10','DD-MON-RR') and
to_date('04-NOV-10','DD-MON-RR')
```

Plan hash value: 909439854

| Id | Operation                   | Name            | Rows | Bytes | Cost (%CPU) | Time     |
|----|-----------------------------|-----------------|------|-------|-------------|----------|
| 0  | SELECT STATEMENT            |                 |      |       | 35 (100)    |          |
| 1  | FILTER                      |                 |      |       |             |          |
| 2  | TABLE ACCESS BY INDEX ROWID | SALES           | 3000 | 166K  | 35 (0)      | 00:00:01 |
| 3  | INDEX RANGE SCAN            | SALES_DATES_NDX | 3000 |       | 9 (0)       | 00:00:01 |

-----

| ----- |                   |       |      |       |             |          |
|-------|-------------------|-------|------|-------|-------------|----------|
| Id    | Operation         | Name  | Rows | Bytes | Cost (%CPU) | Time     |
| ----- |                   |       |      |       |             |          |
| 0     | SELECT STATEMENT  |       |      |       | 244 (100)   |          |
| 1     | FILTER            |       |      |       |             |          |
| 2     | TABLE ACCESS FULL | SALES | 3000 | 166K  | 244 (2)     | 00:00:03 |

We see two different Execution Plans with two different Plan Hash Values.

Why is that so ?

Apparently, the Execution Plan earlier had used an Index Range Scan, using Index “SALES\_DATES\_NDX”.

Later the Execution Plan for the same SQL (because the SQL\_ID is the same) has changed to a Full Table Scan.

What had happened was that I had dropped the Index “SALES\_DATES\_NDX” between the first execution and the second one !

So, this display from AWR shows that the Execution Plans can and have changed over time, with the Plan Hash Value changing accordingly !

Another way is to run \$ORACLE\_HOME/rdbms/admin/awrsqrpt.sql :

Specify the Begin and End Snapshot Ids

~~~~~

Enter value for begin\_snap: 1152

Begin Snapshot Id specified: 1152

Enter value for end\_snap: 1156

End Snapshot Id specified: 1156

Specify the SQL Id

~~~~~

Enter value for sql\_id: 3pat1z2qx4gyg

SQL ID specified: 3pat1z2qx4gyg

Specify the Report Name

~~~~~

The default report file name is awrsqldrpt\_1\_1152\_1156.txt. To use this name, press <return> to continue, otherwise enter an alternative.

Enter value for report\_name:

Using the report name awrsqldrpt\_1\_1152\_1156.txt

# WORKLOAD REPOSITORY SQL Report

## Snapshot Period Summary

| DB Name | DB Id      | Instance | Inst Num | Startup Time    | Release    | RAC |
|---------|------------|----------|----------|-----------------|------------|-----|
| ORCL    | 1229390655 | orcl     | 1        | 21-Apr-11 22:33 | 11.2.0.1.0 | NO  |

|             | Snap Id | Snap Time          | Sessions | Curs/Sess |
|-------------|---------|--------------------|----------|-----------|
| Begin Snap: | 1152    | 21-Apr-11 22:40:02 | 29       | 1.4       |
| End Snap:   | 1156    | 22-Apr-11 00:00:19 | 31       | 1.5       |
| Elapsed:    |         | 80.29 (mins)       |          |           |
| DB Time:    |         | 1.76 (mins)        |          |           |

SQL Summary DB/Inst: ORCL/orcl Snaps: 1152-1156

| SQL Id        | Elapsed Time (ms) |
|---------------|-------------------|
| 3pat1z2qx4gyg | 518               |

Module: sqlplus@localhost.localdomain (TNS V1-V3)  
select sale\_id, cust\_id, remarks from sales where sale\_date between to\_date('01-NOV-10','DD-MON-RR') and to\_date('04-NOV-10','DD-MON-RR')

SQL ID: 3pat1z2qx4gyg

DB/Inst: ORCL/orcl Snaps: 1152-1156

-> 1st Capture and Last Capture Snap IDs

refer to Snapshot IDs within the snapshot range

-> select sale\_id, cust\_id, remarks from sales where sale\_date between to...

| # | Plan Hash<br>Value | Total Elapsed<br>Time (ms) | Executions | 1st Capture<br>Snap ID | Last Capture<br>Snap ID |
|---|--------------------|----------------------------|------------|------------------------|-------------------------|
| 1 | 1231079358         | 400                        | 1          | 1155                   | 1155                    |
| 2 | 909439854          | 118                        | 1          | 1153                   | 1153                    |

This shows that there have been two executions, one of 400ms and the other of 118ms, with two different Plan Hash Values and captured in different Snapshots.

Plan 2 is the “older” plan and took less time. This was the Plan with the Index Range Scan.

Plan 1 (PHV: 1231079358)

-----

Plan Statistics DB/Inst: ORCL/orcl Snaps: 1152-1156  
-> % Total DB Time is the Elapsed Time of the SQL statement divided  
into the Total Database Time multiplied by 100

| Stat Name                  | Statement | Per Execution | % Snap |
|----------------------------|-----------|---------------|--------|
| Elapsed Time (ms)          | 400       | 399.5         | 0.4    |
| CPU Time (ms)              | 248       | 248.0         | 0.3    |
| Executions                 | 1         | N/A           | N/A    |
| Buffer Gets                | 1,281     | 1,281.0       | 0.2    |
| Disk Reads                 | 883       | 883.0         | 4.1    |
| Parse Calls                | 1         | 1.0           | 0.0    |
| Rows                       | 3,000     | 3,000.0       | N/A    |
| User I/O Wait Time (ms)    | 42        | N/A           | N/A    |
| Cluster Wait Time (ms)     | 0         | N/A           | N/A    |
| Application Wait Time (ms) | 0         | N/A           | N/A    |
| Concurrency Wait Time (ms) | 0         | N/A           | N/A    |
| Invalidations              | 0         | N/A           | N/A    |
| Version Count              | 1         | N/A           | N/A    |
| Sharable Mem(KB)           | 14        | N/A           | N/A    |

-----

Execution Plan

| Id | Operation         | Name  | Rows | Bytes | Cost (%CPU) | Time     |
|----|-------------------|-------|------|-------|-------------|----------|
| 0  | SELECT STATEMENT  |       |      |       | 244 (100)   |          |
| 1  | FILTER            |       |      |       |             |          |
| 2  | TABLE ACCESS FULL | SALES | 3000 | 166K  | 244 (2)     | 00:00:03 |

Plan 2 (PHV: 909439854)

-----

Plan Statistics DB/Inst: ORCL/orcl Snaps: 1152-1156  
-> % Total DB Time is the Elapsed Time of the SQL statement divided  
into the Total Database Time multiplied by 100

| Stat Name                  | Statement | Per Execution | % Snap |
|----------------------------|-----------|---------------|--------|
| Elapsed Time (ms)          | 118       | 118.0         | 0.1    |
| CPU Time (ms)              | 45        | 45.0          | 0.1    |
| Executions                 | 1         | N/A           | N/A    |
| Buffer Gets                | 666       | 666.0         | 0.1    |
| Disk Reads                 | 59        | 59.0          | 0.3    |
| Parse Calls                | 1         | 1.0           | 0.0    |
| Rows                       | 3,000     | 3,000.0       | N/A    |
| User I/O Wait Time (ms)    | 84        | N/A           | N/A    |
| Cluster Wait Time (ms)     | 0         | N/A           | N/A    |
| Application Wait Time (ms) | 0         | N/A           | N/A    |
| Concurrency Wait Time (ms) | 0         | N/A           | N/A    |
| Invalidations              | 0         | N/A           | N/A    |
| Version Count              | 1         | N/A           | N/A    |
| Sharable Mem(KB)           | 14        | N/A           | N/A    |

-----

Execution Plan

| Id | Operation                   | Name            | Rows | Bytes | Cost (%CPU) | Time     |
|----|-----------------------------|-----------------|------|-------|-------------|----------|
| 0  | SELECT STATEMENT            |                 |      |       | 35 (100)    |          |
| 1  | FILTER                      |                 |      |       |             |          |
| 2  | TABLE ACCESS BY INDEX ROWID | SALES           | 3000 | 166K  | 35 (0)      | 00:00:01 |
| 3  | INDEX RANGE SCAN            | SALES_DATES_NDX | 3000 |       | 9 (0)       | 00:00:01 |

-----



## Full SQL Text

| SQL ID       | SQL Text                                                                                                                                    |
|--------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| 3pat1z2qx4gy | select sale_id, cust_id, remarks from sales where sale_date between to_date('01-NOV-10', 'DD-MON-RR') and to_date('04-NOV-10', 'DD-MON-RR') |

Report written to awrsqldrpt\_1\_1152\_1156.txt

SQL>