

# Oracle Diagnostics

Hemant K Chitale

# Hemant K Chitale

- whoami ?
- Oracle 5 to Oracle 10gR2 : DOS, Xenix, 8 flavours of Unix, Linux, Windows
- Financial Services, Govt/Not-for-Profit, ERP, Custom
- Production Support, Consulting, Development
- A DBA, not a Developer
- [My Oracle Blog](http://hemantoracledba.blogspot.com) *http://hemantoracledba.blogspot.com*

# Oracle Diagnostics – Joins (1)

- For “Explain Plans” see my earlier presentation [OracleDiagnostics\\_Explain\\_Plans\\_simple](#) (or at WizIQ)
- This presentation is based on the case study [Nested Loops and Consistent Gets](#)

# A Nested Loop join

- In a Nested Loop join, Oracle uses the rowset retrieved from the “outer” (driving) table to query the “inner” (driven) table
- The Optimizer would choose a Nested Loop join if it expects to be able to retrieve the rowset of the outer table very quickly \*and\* expects that the size would be small enough to not require querying the inner table too frequently

Given this query :

```
SQL> explain plan for
  2  select  product_desc, txn_id, transaction_amt
  3  from transactions t, product_table p
  4  where txn_date between to_date('01-FEB-2011','DD-MON-YYYY')
and to_date('28-FEB-2011','DD-MON-YYYY')
  5  and txn_id between 155000 and 156000
  6  and country_cd = 'IN'
  7  and t.product_cd=p.product_cd
  8  /
```

We expect a Join between TRANSACTIONS and PRODUCT\_TABLE

```
SQL> select * from table(dbms_xplan.display);
```

```
PLAN_TABLE_OUTPUT
```

```
-----  
-----  
Plan hash value: 1299935411  
  
-----  
-----
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
Pstart	Pstop						
0	SELECT STATEMENT		1	63	25 (0)	00:00:01	
1	NESTED LOOPS						
2	NESTED LOOPS		1	63	25 (0)	00:00:01	
3	PARTITION RANGE SINGLE		1	51	24 (0)	00:00:01	
2   2							
* 4	TABLE ACCESS BY LOCAL INDEX ROWID	TRANSACTIONS	1	51	24 (0)	00:00:01	
2   2							
* 5	INDEX RANGE SCAN	TRANSACTIONS_NDX	252		4 (0)	00:00:01	
2   2							
* 6	INDEX UNIQUE SCAN	SYS_C0016611	1		0 (0)	00:00:01	
7	TABLE ACCESS BY INDEX ROWID	PRODUCT_TABLE	1	12	1 (0)	00:00:01	

```
-----  
Predicate Information (identified by operation id):  
-----
```

```
4 - filter("TXN_DATE"<=TO_DATE(' 2011-02-28 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))  
5 - access("COUNTRY_CD"='IN' AND "TXN_ID">=155000 AND "TXN_ID"<=156000)  
6 - access("T"."PRODUCT_CD"="P"."PRODUCT_CD")
```

```
SQL>
```

Here :

	3		PARTITION RANGE SINGLE				1		51		24	(0)		00:00:01
	2		2											
*	4		TABLE ACCESS BY LOCAL INDEX ROWID		TRANSACTIONS		1		51		24	(0)		00:00:01
	2		2											
*	5		INDEX RANGE SCAN		TRANSACTIONS_NDX		252				4	(0)		00:00:01
	2		2											

Predicate Information :

```
4 - filter("TXN_DATE"<=TO_DATE(' 2011-02-28 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
5 - access("COUNTRY_CD"='IN' AND "TXN_ID">=155000 AND "TXN_ID"<=156000)
```

The execution is done in this sequence :

**Step 1 : Operation 5 : Index Range Scan** (expected to retrieve 252 rowids),  
accessing the TRANSACTIONS\_NDX index by COUNTRY\_CD = 'IN' and  
TXN\_ID between 155000 and 156000

**Step 2 : Operation 4 : Table Access of TRANSACTIONS**, using the RowIDs  
fetched from Operation 5. For every row fetched by RowID, filter for  
TXN\_DATE less than 28-Feb-11. Expect to return only 1 single row after  
the filter for TXN\_DATE (note : This is key to the optimizer's choice of a Nested  
Loop join)

**Parent Step 3 : Operation 3 : Determine** that this is a Range Scan within  
Partition 2 (Start and Stop Partition IDs being "2" and "2"). This range scan will  
return the single row from Operation 4.

## Continuing with the other operations :

	2		NESTED LOOPS				1		63		25	(0)		00:00:01	
	3		PARTITION RANGE SINGLE				1		51		24	(0)		00:00:01	
	2		2												
	*	6		INDEX UNIQUE SCAN		SYS_C0016611		1				0	(0)		00:00:01

Predicate Information :

6 - access("T"."PRODUCT\_CD"="P"."PRODUCT\_CD")

**Step 4 Operation 2 :** For the expected single row returned from the Partition Range (in Operation 3), call Operation 6 to scan the Unique Index SYS\_C0016611. The Index Unique Scan is expected to be executed once only (as Operation 3 returns 1 row) and, in-turn, expected to return a single RowID (being a Unique Index scan)

This Index is the Primary Key index for PRODUCT\_TABLE. The Index Scan is for the PRODUCT\_CD retrieved the previous step.

	1		NESTED LOOPS											
	7		TABLE ACCESS BY INDEX ROWID		PRODUCT_TABLE		1		12		1	(0)		00:00:01

**Step 6 Operation 1 :** For the single RowID expected from Operation 6, access the PRODUCT\_TABLE. (Since the previous step expects to return a single RowID, this step, too, expects to read a single row from the table).



# Hash Join

- This is a demonstration of a Hash Join from the same case study
- In a Hash Join, Oracle attempts to retrieve all the join column values from the first table, compute “hash” values and place them in memory
- It then expects to probe the second table and compare the values with the in-memory hash table. A match is a “join”
- Performance of a Hash Join depends on the size of the join key and the available memory for the in-memory Hash Table

For this query :

```
SQL> explain plan for
  2  select /*+ USE_HASH (t p) */ product_desc, txn_id,
transaction_amt
  3  from transactions t, product_table p
  4  where txn_date between to_date('01-FEB-2011','DD-MON-YYYY')
and to_date('28-FEB-2011','DD-MON-YYYY')
  5  and txn_id between 155000 and 156000
  6  and country_cd = 'IN'
  7  and t.product_cd=p.product_cd
  8  /
```

We have explicitly provided a Hint (directive) to Oracle to execute a Hash Join between the two tables

# It is

```
SQL> select * from table(dbms_xplan.display);
```

```
PLAN_TABLE_OUTPUT
```

```
-----  
Plan hash value: 971735053  
  
-----
```

```
-----  
| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |  
Pstart| Pstop |  
-----  
| 0 | SELECT STATEMENT | | 1 | 63 | 28 (4) | 00:00:01 |  
| | | | | | | |  
|* 1 | HASH JOIN | | 1 | 63 | 28 (4) | 00:00:01 |  
| | | | | | | |  
| 2 | PARTITION RANGE SINGLE | | 1 | 51 | 24 (0) | 00:00:01 |  
2 | 2 | | | | | | |  
|* 3 | TABLE ACCESS BY LOCAL INDEX ROWID | TRANSACTIONS | 1 | 51 | 24 (0) | 00:00:01 |  
2 | 2 | | | | | | |  
|* 4 | INDEX RANGE SCAN | TRANSACTIONS_NDX | 252 | | 4 (0) | 00:00:01 |  
2 | 2 | | | | | | |  
| 5 | TABLE ACCESS FULL | PRODUCT_TABLE | 14 | 168 | 3 (0) | 00:00:01 |  
| | | | | | | |  
-----
```

```
-----  
Predicate Information (identified by operation id):  
-----
```

- 1 - access("T"."PRODUCT\_CD"="P"."PRODUCT\_CD")
- 3 - filter("TXN\_DATE"<=TO\_DATE(' 2011-02-28 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
- 4 - access("COUNTRY\_CD"='IN' AND "TXN\_ID">=155000 AND "TXN\_ID"<=156000)

```
SQL>
```

Here :

	2		PARTITION RANGE SINGLE				1		51		24		(0)		00:00:01	
2		2														
*	3		TABLE ACCESS BY LOCAL INDEX ROWID		TRANSACTIONS		1		51		24		(0)		00:00:01	
2		2														
*	4		INDEX RANGE SCAN		TRANSACTIONS_NDX		252				4		(0)		00:00:01	
2		2														

Predicate Information :

3 - filter("TXN\_DATE"<=TO\_DATE(' 2011-02-28 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))

4 - access("COUNTRY\_CD"='IN' AND "TXN\_ID">=155000 AND "TXN\_ID"<=156000)

**Step 1 : Operation 4 : Retrieve 252 RowIDs from the TRANSACTIONS\_NDX**

**Step 2 : Operation 3 : Use the 252 RowIDs to lookup the TRANSACTIONS table and filter to TXN\_DATE. Expect to return a single row**

**Parent Step 3 : Confirm that these steps are executed against a single Partition**

Here :

	*	1		HASH JOIN				1		63		28		(4)		00:00:01	
		2		PARTITION RANGE SINGLE				1		51		24		(0)		00:00:01	
2			2														

		5		TABLE ACCESS FULL		PRODUCT_TABLE		14		168		3		(0)		00:00:01	

Predicate Information :

1 - access("T"."PRODUCT\_CD"="P"."PRODUCT\_CD")

**Step 4 Operation 1:** For the single row returned from the Partition Range Scan, execute a Hash Join. The hashed value(s) for this row are stored in an in-memory hash area

**Step 5 : Operation 5 :** Probe the `PRODUCT_TABLE` (via a FullTableScan) to read 14 rows. Place the values in memory and attempt a join on `PRODUCT_CD`.