

Partitioning Tables and Indexing Them

Hemant K Chitale

Product Specialist, Standard Chartered Bank
Oracle ACE

© Hemant K Chitale <http://hemantoracledba.blogspot.com>

`whoami`

- DBA with 20 years experience on wide variety of platforms
- DBA team lead and consultant
- Financial, Manufacturing, Government, not-for-profit
- Mission critical, reporting and “nice-to-have” databases

Introduction

- Very Large Databases
- Growth of Transaction Volumes
- Need to retain Historical Data
- Performance Issues from querying large tables
- Separation of OLTP and DWH/DSS systems
- *Question : Is Partitioning only for Very Large Tables ?*
- What you might not find on Google

Pre-Oracle 8

- Partitioning was introduced in V8
- UNION-ALL joins of distinct tables encapsulated in Views
- You have to control which table a new record goes into
- Check Constraints in V7.3

Elements

- Each Partition has the same Logical Attributes as the Table :
 - Column Names, Column Ordering, Datatypes
 - Constraints
- Partitions are distinct Segments. Therefore, they have distinct Physical Attributes :
 - PCTFREE
 - COMPRESSION
 - Tablespace

Elements – Example

```
SQL> create table SALES_TABLE(sale_date date not null, region varchar2(8), sale_qty
    number)
    2  partition by range (sale_date) subpartition by list (region)
    3  (
    4  partition p_2010 values less than (to_date('01-JAN-2011','DD-MON-YYYY'))
    5      (subpartition p_2010_s_east values ('EAST'),
    6          subpartition p_2010_s_north values ('NORTH'),
    7          subpartition p_2010_s_south values ('SOUTH'),
    8          subpartition p_2010_s_west values ('WEST')
    9      )
   10  ,
   11  partition p_2011 values less than (to_date('01-JAN-2012','DD-MON-YYYY'))
   12      (subpartition p_2011_s_east values ('EAST'),
   13          subpartition p_2011_s_north values ('NORTH'),
   14          subpartition p_2011_s_south values ('SOUTH'),
   15          subpartition p_2011_s_west values ('WEST')
   16      )
   17  )
   18  /
```

Table created.

SQL>

Elements – Example (2)

```
SQL> select object_id, object_name, subobject_name, object_type
  2   from user_objects
  3   order by object_type, object_name, subobject_name
  4   /
```

OBJECT_ID	OBJECT_NAME	SUBOBJECT_NAME	OBJECT_TYPE
54889	SALES_TABLE		TABLE
54890	SALES_TABLE	P_2010	TABLE PARTITION
54891	SALES_TABLE	P_2011	TABLE PARTITION
54892	SALES_TABLE	P_2010_S_EAST	TABLE SUBPARTITION
54893	SALES_TABLE	P_2010_S_NORTH	TABLE SUBPARTITION
54894	SALES_TABLE	P_2010_S_SOUTH	TABLE SUBPARTITION
54895	SALES_TABLE	P_2010_S_WEST	TABLE SUBPARTITION
54896	SALES_TABLE	P_2011_S_EAST	TABLE SUBPARTITION
54897	SALES_TABLE	P_2011_S_NORTH	TABLE SUBPARTITION
54898	SALES_TABLE	P_2011_S_SOUTH	TABLE SUBPARTITION
54899	SALES_TABLE	P_2011_S_WEST	TABLE SUBPARTITION

11 rows selected.

```
SQL>
```

Elements – Example (3)

```
SQL> select segment_name, partition_name, segment_type, tablespace_name
  2   from user_segments
  3   order by segment_name, partition_name
  4   /
```

SEGMENT_NAME	PARTITION_NAME	SEGMENT_TYPE	TABLESPACE_NAME
SALES_TABLE	P_2010_S_EAST	TABLE SUBPARTITION	USERS
SALES_TABLE	P_2010_S_NORTH	TABLE SUBPARTITION	USERS
SALES_TABLE	P_2010_S_SOUTH	TABLE SUBPARTITION	USERS
SALES_TABLE	P_2010_S_WEST	TABLE SUBPARTITION	USERS
SALES_TABLE	P_2011_S_EAST	TABLE SUBPARTITION	USERS
SALES_TABLE	P_2011_S_NORTH	TABLE SUBPARTITION	USERS
SALES_TABLE	P_2011_S_SOUTH	TABLE SUBPARTITION	USERS
SALES_TABLE	P_2011_S_WEST	TABLE SUBPARTITION	USERS

8 rows selected.

```
SQL>
```


Elements – Example (4)

```
SQL> select partition_name, partition_position, high_value
       2  from user_tab_partitions
       3  where table_name = 'SALES_TABLE'
       4  order by partition_position
       5  /
```

PARTITION_NAME	PARTITION_POSITION	HIGH_VALUE
P_2010	1	TO_DATE(' 2011-01-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS', 'NLS_CALENDAR=GREGORIA
P_2011	2	TO_DATE(' 2012-01-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

```
SQL>
```

Elements – Example (5)

```
SQL> select partition_name, subpartition_name, subpartition_position, high_value
      2  from user_tab_subpartitions
      3  where table_name = 'SALES_TABLE'
      4  order by partition_name, subpartition_position
      5  /
```

PARTITION_NAME	SUBPARTITION_NAME	SUBPARTITION_POSITION
----------------	-------------------	-----------------------

HIGH_VALUE

P_2010	P_2010_S_EAST	1
'EAST'		
P_2010	P_2010_S_NORTH	2
'NORTH'		
P_2010	P_2010_S_SOUTH	3
'SOUTH'		
P_2010	P_2010_S_WEST	4
'WEST'		
P_2011	P_2011_S_EAST	1
'EAST'		
P_2011	P_2011_S_NORTH	2
'NORTH'		
P_2011	P_2011_S_SOUTH	3
'SOUTH'		
P_2011	P_2011_S_WEST	4
'WEST'		

SQL>

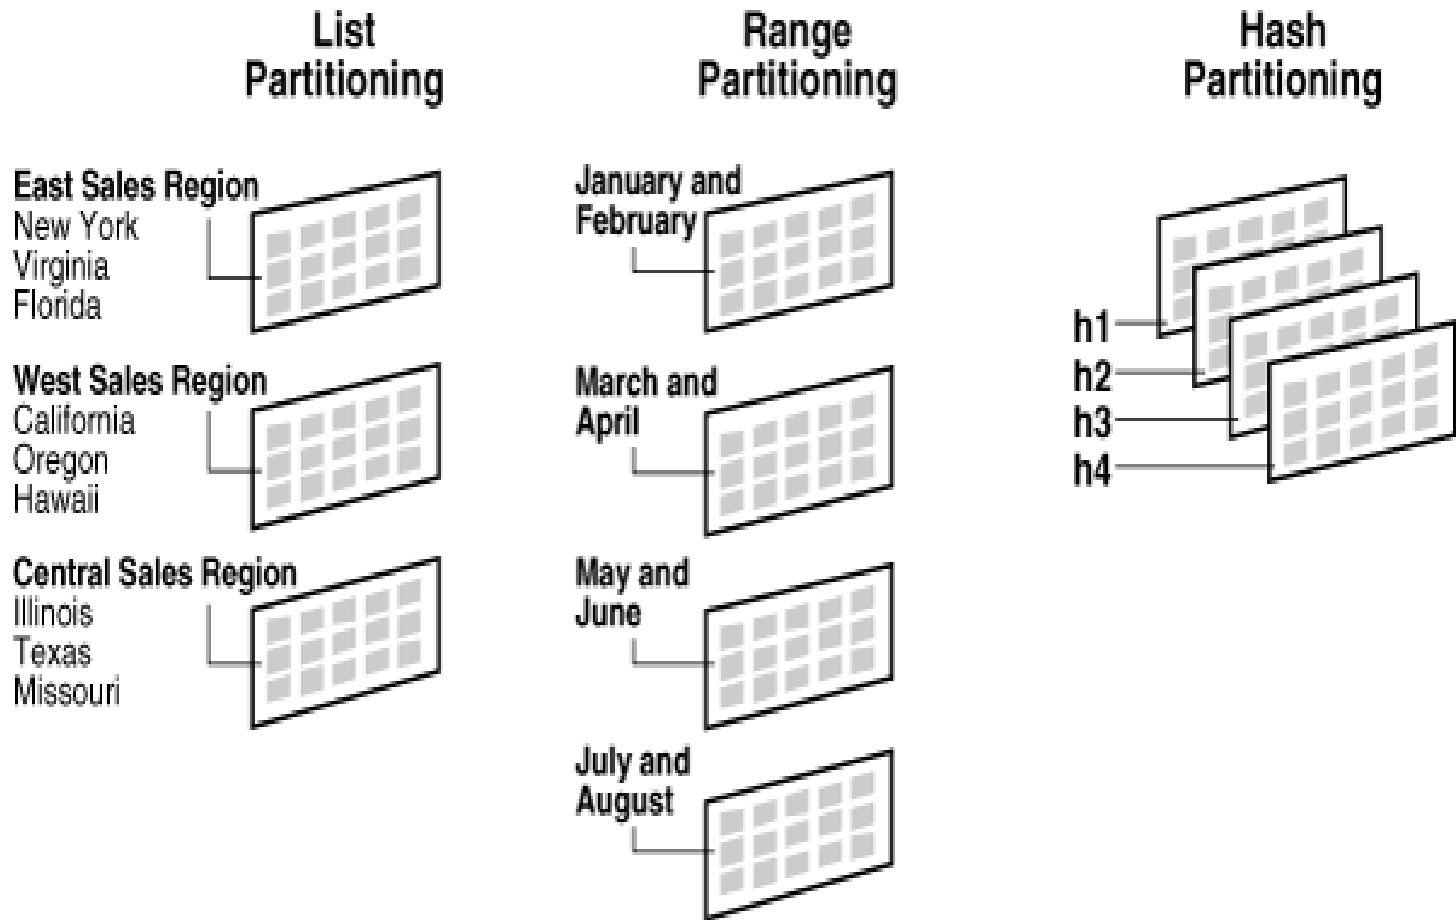
Common Types

- Range Partitioning (*introduced in V8 and still the most popular*)
- Hash Partitioning (*introduced in 8i but much less used than Range and List*)
- List Partitioning (*introduced in 9i*)
- Composite (Range-Hash) (Range-List) (List-Range) (List-Hash) etc Partitioning

Complex Definitions (11g)

- Virtual Column Partitioning
- Reference Partitioning
- Interval Partitioning (as an extension to Range Partitioning)
- *** not in scope ***

Partitioning Types



From the 11gR2 VLDB and Partitioning Guide : Figure 2-2

Range Partitioning

- Most frequently used with Date Ranges

```
partition by range (sale_date)
(partition p_2010 values less than
(to_date('01-JAN-2011','DD-MON-YYYY')))
```

- Useful when you need to be able to Archive/Purge by Date (simply TRUNCATE/DROP the oldest Partition(s))
- Supports multi-column Partition Key
- Each Partition stores values within it's Upper Bound
- Use a MAXVALUE partition for values above known max
- The Optimizer can use the min and max values for each Partition

Range Partitioning – Examaple (1)

```
SQL> create table ACCOUNTING
 2  (biz_country varchar2(10) not null, acctg_year number not
 3  null, data_1 varchar2(20))
 4  partition by range (biz_country, acctg_year)
 5  (
 6  partition p_in_2006 values less than ('IN',2007),
 7  partition p_in_2007 values less than ('IN',2008),
 8  partition p_in_2008 values less than ('IN',2009),
 9  partition p_sg_2006 values less than ('SG',2007),
10  partition p_sg_2007 values less than ('SG',2008),
11  partition p_sg_2008 values less than ('SG',2009),
12  partition p_max values less than (MAXVALUE, MAXVALUE)
13  )
14  /
```

Table created.

```
SQL>
```

Range Partitioning – Example (2)

```
SQL> insert into ACCOUNTING values ('IN',2007,'Row 1');
```

```
1 row created.
```

```
SQL> insert into ACCOUNTING values ('IN',2008,'Row 2');
```

```
1 row created.
```

```
SQL> insert into ACCOUNTING values ('JP',2007,'Row 3');
```

```
1 row created.
```

```
SQL> insert into ACCOUNTING values ('JP',2015,'Row 4');
```

```
1 row created.
```

```
SQL> insert into ACCOUNTING values ('US',2006,'Row 5');
```

```
1 row created.
```

```
SQL> insert into ACCOUNTING values ('US',2009,'Row 6');
```

```
1 row created.
```

```
SQL>
```


Range Partitioning – Example (3)

```
SQL> select * from ACCOUNTING partition (p_in_2006);  
no rows selected
```

```
SQL> select * from ACCOUNTING partition (p_in_2007);  
BIZ_COUNTR ACCTG_YEAR DATA_1  
-----  
IN          2007 Row 1
```

```
SQL> select * from ACCOUNTING partition (p_in_2008);  
BIZ_COUNTR ACCTG_YEAR DATA_1  
-----  
IN          2008 Row 2
```

```
SQL> select * from ACCOUNTING partition (p_sg_2006);  
BIZ_COUNTR ACCTG_YEAR DATA_1  
-----  
JP          2007 Row 3  
JP          2015 Row 4
```

```
SQL> select * from ACCOUNTING partition (p_max);  
BIZ_COUNTR ACCTG_YEAR DATA_1  
-----  
US          2006 Row 5  
US          2009 Row 6
```

```
SQL>
```

Hash Partitioning

- Define the Partition Key
- Oracle dynamically uses the PK Value to allocate a row to a Partition – you cannot associate them in advance
- A Hashing algorithm is used
- Useful when you have a large number of values but cannot determine allocation
- Define 2^N Partitions else allocation is unbalanced

List Partitioning

- Well-defined list of values for the Partition Key
- Single Column only
- (11g allows List-List composite partitioning)
- Use a DEFAULT Partition for unknown values
- The Optimizer *knows* that every row has the same value for the Partition Key

List Partitioning – Example (1)

- Badly defined Range Partitioning :

```
SQL> create table MONTH_END_BALANCES
  2  (Partition_Key      varchar2(8) not null, account_number number, balance number)
  3  partition by Range (Partition_Key)
  4  (partition P_2011_JAN values less than ('20110132'),
  5  partition P_2011_FEB values less than ('20110229'),
  6  partition P_2011_MAR values less than ('20110332'),
  7  partition P_2011_APR values less than ('20110431'),
  8  partition P_2011_MAY values less than ('20110532'),
  9  partition P_2011_JUN values less than ('20110631'),
 10  partition P_2011_JUL values less than ('20110732'),
 11  partition P_2011_AUG values less than ('20110832'),
 12  partition P_2011_SEP values less than ('20110931'),
 13  partition P_2011_OCT values less than ('20111032'),
 14  partition P_2011_NOV values less than ('20111131'),
 15  partition P_2011_DEC values less than ('20111232')
 16  )
 17  /
```

Table created.

SQL>

List Partitioning – Example (2)

- Better Definition :


```
SQL> create table MONTH_END_BALANCES
  2  (Partition_Key      varchar2(6) not null, account_number number, balance number)
  3  partition by List (Partition_Key)
  4  (partition P_2011_JAN values ('201101'),
  5  partition P_2011_FEB values ('201102'),
  6  partition P_2011_MAR values ('201103'),
  7  partition P_2011_APR values ('201104'),
  8  partition P_2011_MAY values ('201105'),
  9  partition P_2011_JUN values ('201106'),
 10  partition P_2011_JUL values ('201107'),
 11  partition P_2011_AUG values ('201108'),
 12  partition P_2011_SEP values ('201109'),
 13  partition P_2011_OCT values ('201110'),
 14  partition P_2011_NOV values ('201111'),
 15  partition P_2011_DEC values ('201112')
 16  )
 17  /
```

Table created.

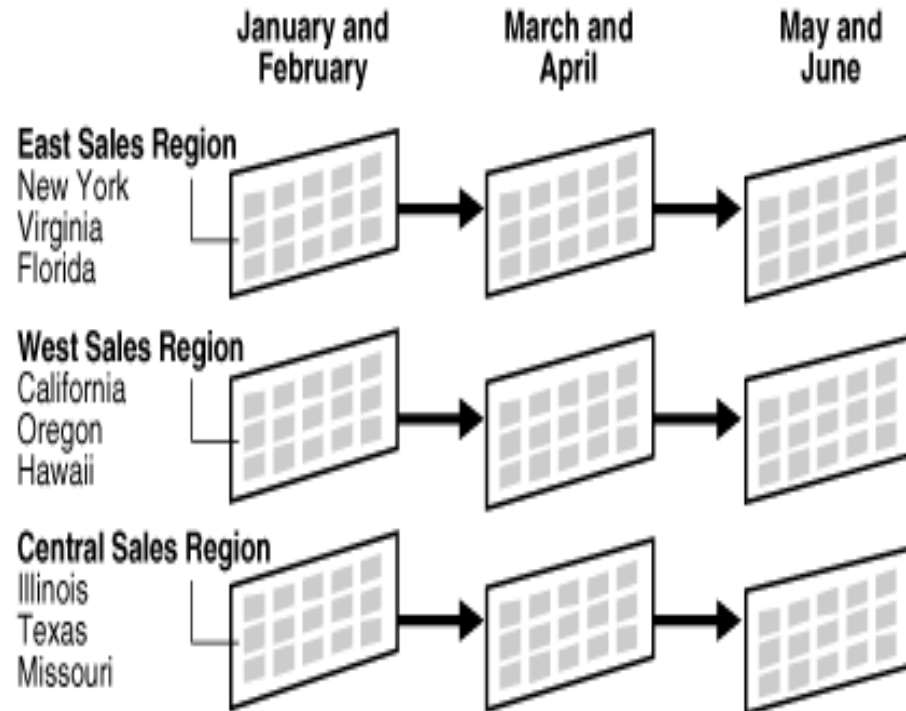
SQL>

Composite Partitioning

Composite Partitioning
Range-Hash



Composite Partitioning
Range - List



From the 11gR2 VLDB and Partitioning Guide Figure 2-3

Choosing a Partitioning Method

- Need to Archive / Purge Data : Range
- Querying only for specific range (month in Range Partition, column value in Hash or List)
- Distribution of large data set : Hash
- Discrete, small, set of values : List

Partitioning Method – Examples

- Historical data of Transactions that needs to be purged after 7years : DATE Range
- Employee / Contractor information by State : STATE List
- Statistical Information with a large range of values : VALUE Hash

Adding Data

- You do not need to specify the target Partition name when running an INSERT (serial, parallel, direct).
- If you do name the target Partition, Oracle will still check if the data belongs to the Partition.
- Bulk Insert (Direct Path) into a single (named) Partition will lock only that Partition.
- Use EXCHANGE Partition to switch a non-Partitioned Table with a Partition

Maintaining Partitioned Tables

- Operations : ADD, DROP/TRUNCATE, COALESCE, SPLIT, MERGE, EXCHANGE, DBMS_REDEFINITION, MOVE
- However there are caveats and restrictions
- Maintenance has impact on Indexes as Maintenance Operations are DDLs

ADD Partition

- ADD is to create a new Partition.
- In Range Partitioning you cannot “add” an intermediate partition, you have to SPLIT a Partition (as also if you have a MAXVALUE Partition)
- In Hash Partitioning, adding a new Partition results in Oracle actually splitting an existing Partition – ending up with unbalanced Partitions

DROP or TRUNCATE

- DROP is to drop a Partition
- In Range Partitioning, new rows will go into the “next” Partition
- In List Partitioning, reinserting the dropped values requires a DEFAULT Partition
- TRUNCATE truncates a Partition (TRUNCATE Table truncates all the Partitions) but retains the definition

COALESCE, SPLIT and MERGE

- MERGE allows you to merge adjacent Partitions (use COALESCE in Hash Partitioning to reduce one of the partitions)
- Rows in the “removed” Partition are “moved” into the “new” Partition
- MAXVALUE and DEFAULT Partitions can be SPLITted to create “new” Partitions
- Existing Range Partitions can be SPLITted for granularity

EXCHANGE

- EXCHANGE allows you to interchange a non-Partitioned Table with a Partition
- The empty object is replaced – so be careful as to which of the two is empty !
- EXCHANGE is used to load data from a staging table without another INSERT operation
- EXCHANGE is useful to “move out” an older Partition as a separate Table that can be Exported / Archived

DBMS_REDEFINITION and MOVE

- A Partition can be moved to another Tablespace via DBMS_REDEFINITION
- A Partition can be MOVEd in the same manner as an ALTER TABLE ... MOVE
- DBMS_REDEFINITION can also be used “convert” a non-Partitioned Table to a Partitioned Table by copying data across online

Maintaining Indexes

- Types of Indexes
 - Global
 - Global Partitioned
 - Local
- Local Indexes are Equi-Partitioned with the Table
- Global Indexes become UNUSABLE with Table Partition Maintenance unless they are UPDATED

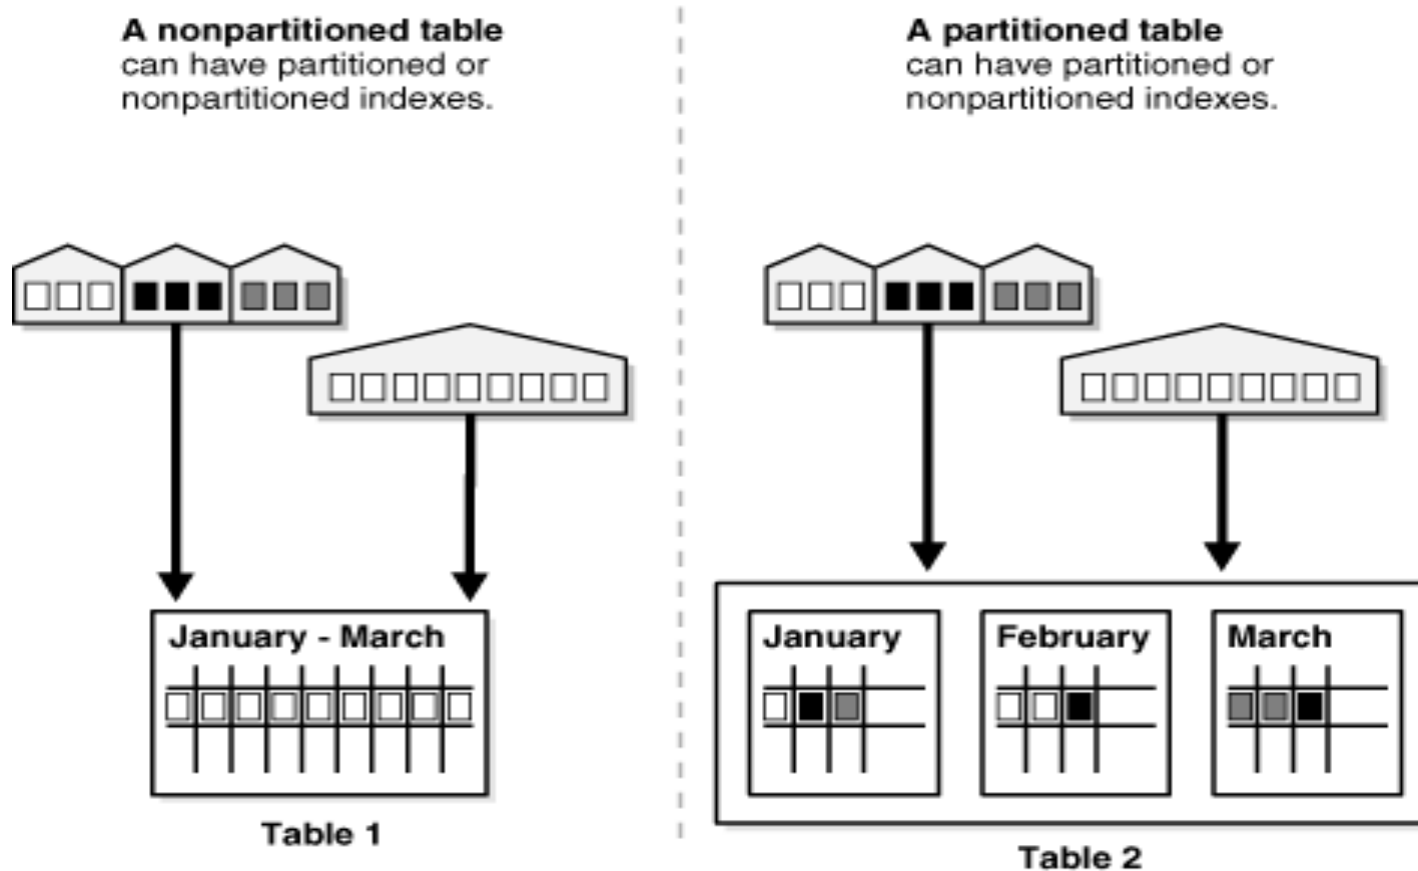
Global Indexes

- Generally recommended in “OLTP” environments --- the real rationale is that queries do not do Partition Pruning
- No different from Local Indexes when retrieving a single row but much more useful when retrieving rows across Partitions
- Global Partitioned Indexes are Partitioned on a separate Partition Key. Useful for queries that are Index-only

Local Indexes


- Generally used for DWH/DSS ; queries that target Partitions (Partition Pruning)
- Must contain Partition Key if defined as UNIQUE
- Most Partition Maintenance operations can also maintain Index Partitions
 - Adding a Table Partition adds Index Partition
 - Dropping a Table Partition drops Index Partition

Partitioned Indexes




From the 11gR2 VLDB and Partitioning Guide : Figure 2-1

Global Indexes



From the 11gR2 VLDB and Partitioning Guide Figure 2-8

Local Indexes



Performance Strategies

- Partition Pruning for subset(s) of Table
- Bulk Insert with Parallel (across Partitions)
- Global Indexes for non-partition Pruning
- Partition-Wise Join with Equi-Partitioning on the same Partition Key
- Hash Partitions for Partition-Wise Joins
- Verify Index definitions when attempting EXCHANGE

Archiving Data

- Data Partitioned on Date key can be Archived/Purged with EXCHANGE and/or DROP
- “Older” Partitions can be moved to slower storage with MOVE
- Partitions can be SPLITted or MERGEEd as required

Common Mistakes

- Using the wrong Partitioning Method (and using the wrong datatype for the Partition Key)
- Incorrect Ranges, data going into “other” Partitions (caution : TRUNCATE or DELETE ... PARTITION can result in the wrong data being purged !)
- Unequal Partitions (Range or List or Hash)

Common Mistakes - 2

- Updating the Partition Key
- Incorrect LOCAL indexes --- which do not get used or suffer overuse --- where GLOBAL indexes would have helped
- Partition Maintenance with GLOBAL indexes present but not UPDATEing them

Conclusion

- ***DON'T*** rush into Partition
- Carefully consider :
 - Manner in which data is being or is planned to be inserted
 - Queries
 - Need or plans to purge / archive data
 - Queries
 - Index definitions
 - Queries

Thank you !

- Visit my Oracle Blog
- <http://hemantoracledba.blogspot.com>