

SQL in Functions

This article was submitted for AIOUG's CONNECT Journal in November 2010

Author: Hemant K Chitale Oracle ACE.

About: Hemant K Chitale has more than 20 years Oracle DBA experience on a large number of platforms. He has been in production support and consulting roles and has worked for or with organisations in financial services, manufacturing, Government and not-for-profit sectors. He has provided Oracle DBA support for a wide variety of applications.

Introduction:

As DBAs and Developers, we are familiar with both SQL and PLSQL. Where the two "languages" meet is when SQL is used in PLSQL Blocks – as in Procedures, Triggers and Functions. The usage of SQL in custom Functions is a special case. ("custom Functions" are our own developed code as opposed to Oracle supplied Functions that are part of the kernel – examples are TO_DATE / TO_CHAR, DECODE, NVL). There are two issues that arise:

- a. The SQL is executed each time the Function is called – thus including a Function call for each row in a "parent" SQL results in repeated executions of the Function's SQL
- b. Oracle's Read Consistency model, being at the *statement* level is applied to each execution of the SQL separately. This means that data can "seem" to be inconsistent if it is simultaneously updated from another session while it is being retrieved repeatedly by a Function. I plan to demonstrate this in a subsequent article

Description:

In this paper, I walk through an example of the first type of issue with SQL in Functions. (I plan to demonstrate the second issue in a subsequent paper). I demonstrate a case where a PLSQL Function (containing an SQL statement) is used inside an SQL that operates on many rows. I show how the recursive calls have an impact on the total "effort" that Oracle undertakes. Repeated switching between PLSQL and SQL also has significant CPU overheads – "context switches" have to take place, although Oracle has merged the SQL and PLSQL engines considerably, beginning with 9i. As it is PLSQL code that loops "for each row" is a performance killer. Adding SQL recursive calls into the loop doesn't help.

Development Practices:

Common Development "Best Practices" generally include the specification of "Use ReUsable Code". In the context of Oracle this means Packages, Stored Procedures and Functions which can be called from various points in the overall Application. ReUsable Code is perfectly fine where it helps avoid

- a. Overheads in Parse Time
- b. Latch Contention and Space Management in the Shared Pool
- c. Ease of Development (avoid re-writing the same logic repeatedly, in different sections of the application and, possibly, even share the same code across multiple applications).

Thus, it is easy to write a Function and then invoke the same Function wherever it is necessary in the application.

Case:

In this Test Case, I create:

Tables:

FX_RATES: This stores foreign exchange (“forex”) rates for 26 countries over the past 1000 days. The Currency Code and Conversion Rate Date together form the Primary Key.

TRANSACTIONS: This has 100,000 Transactions in the various currencies. The last two columns of this table "CONVERTED_AMT" and "UPDATED_AT" are NULL when the transactions are created. The Transaction ID (an increasing Sequence) is the Primary Key.

Function:

CONVERT_FX: This Function returns the Converted Amount for a given Transaction Amount based on the Transaction Date and Currency Code. It does a lookup on the FX Rates table to do so.

This function makes perfect sense when used in a Forms / Screen / Browser based application where a converted amount has to be presented to a user who is viewing 1 or 5 or 10 records only. Thus, multiple screens can invoke the same function and it does not need to be rewritten into each screen.

User Requirements:

The CONVERTED_AMT and UPDATED_AT columns in the TRANSACTIONS table must be populated by a Batch Program (possibly running weekly or monthly). This may later be used for reporting the summary of Debits and Credits.

The Design of the Application and Batch Program:

By design, the Function, as a stand-alone, is sound. The CONVERT_FX Function is a quick, seemingly, lightweight operation as the FX_RATES table is a small table and a lookup for any particular Currency Code + Conversion Rate Date is very fast. However, the issue in this Test Case is when the Lead Developer decides to use this Function in his Batch Program – in a large UPDATE. The Function is Row oriented – it can be invoked only for a particular combination of Amount + Date + Currency. The Function is then used in an UPDATE statement against the TRANSACTIONS table. Since the Function has to be executed against each row, the UPDATE, unfortunately, ends up invoking the SQL within the Function repeatedly -- once for each row that it is to be executed upon in the Transactions table. For simplicity, I have assumed that the Batch Job has to run against all the 100,000 rows in the Transactions Table.

Data Creation:

This is the creation of the data for this Test Case:

```
SQL>
SQL> -- create the FX_RATES table
SQL> drop table fx_rates;
```

Table dropped.

```
SQL> create table fx_rates
  2  ( currency_code varchar2(1) not null,
  3    conv_rate_date date not null,
  4    conv_rate number(9,4))
  5  /
```

Table created.

```
SQL> alter table fx_rates add constraint my_fx_rates_pk primary key
(currency_code, conv_rate_date);
```

Table altered.

```
SQL>
SQL>
SQL> -- create the TRANSACTIONS table
SQL> drop table transactions;
```

Table dropped.

```
SQL> create table transactions
  2  (transaction_id number not null,
  3    transaction_type_code varchar2(1),
  4    transaction_date date,
  5    transaction_amount number,
  6    transaction_currency_code varchar2(1),
  7    transactions_reference_text varchar2(250),
  8    converted_amount number,
  9    updated_at date )
 10  /
```

Table created.

```
SQL> alter table transactions add constraint my_transactions_pk primary key
(transaction_id);
```

Table altered.

```
SQL>
SQL>
SQL> -- populate the FX_RATES table
SQL> variable i number;
SQL> variable j number;
SQL> variable c number;
SQL>
SQL> begin
  2  for i in 1..1000
  3  loop
  4    for c in 1..26
  5  loop
  6    insert into fx_rates values
```

(c) Hemant K Chitale 2010

SQL in Functions

```
7      (  
8      chr(c+64), -- 26 different Currency Codes  
9      trunc(sysdate)-i, -- 1000 days  
10     trunc((dbms_random.value(4,100)),3) -- Conversion Rates  
11     ) ;  
12 end loop;  
13 end loop;  
14 end ;  
15 /
```

PL/SQL procedure successfully completed.

```
SQL>  
SQL>  
SQL> -- populate the TRANSACTIONS table  
SQL> begin  
2   for j in 1..100000  
3   loop  
4   insert into transactions values  
5   (  
6   j,  
7   decode(mod(j,2),0,'D',1,'C'), -- Transaction Type  
8   trunc(sysdate)-(j*1000/100000), -- days  
9   trunc(dbms_random.value(5,5000),2), -- Transaction Amounts  
10  dbms_random.string('U',1), -- One of 26 Currency Codes  
11  rpad(dbms_random.string('a',25),200,'X'), -- Txn Reference (text)  
12  NULL, -- Converted Amount  
13  NULL -- Update date  
14  ) ;  
15  end loop;  
16 end;  
17 /
```

PL/SQL procedure successfully completed.

```
SQL>  
SQL> commit;
```

Commit complete.

```
SQL>  
SQL> exec dbms_stats.gather_table_stats(user,'FX_RATES',method_opt=>'FOR  
ALL COLUMNS SIZE 1',estimate_percent=>100,cascade=>TRUE);
```

PL/SQL procedure successfully completed.

```
SQL> exec  
dbms_stats.gather_table_stats(user,'TRANSACTIONS',method_opt=>'FOR ALL  
COLUMNS SIZE 1',estimate_percent=>100,cascade=>TRUE);
```

PL/SQL procedure successfully completed.

```
SQL>  
SQL> select count(*) from fx_rates;
```

```
   COUNT(*)  
-----  
      26000
```

SQL in Functions

```
SQL> select count(distinct(currency_code)) from fx_rates;
```

```
COUNT(DISTINCT(CURRENCY_CODE))
-----
                             26
```

```
SQL> select count(*) from transactions;
```

```
COUNT(*)
-----
    100000
```

```
SQL>
```

```
SQL> REM fx_rates has rates for 26 currencies and 1000 days
```

```
SQL> REM transactions has 100,000 transactions for different currencies
```

```
SQL>
```

```
SQL>
```

```
SQL> -- create the function to convert transaction amount
```

```
SQL> create or replace function
```

```
    2 convert_fx (currency_code_in in varchar2, amount_in in number,
txn_date_in in date)
```

```
    3 return number
```

```
    4 is
```

```
    5 converted_amt number;
```

```
    6 begin
```

```
    7     select amount_in*conv_rate into converted_amt from fx_rates where
currency_code=currency_code_in and conv_rate_date=txn_date_in ;
```

```
    8     return converted_amt;
```

```
    9 end;
```

```
   10 /
```

Function created.

```
SQL> show errors
```

No errors.

```
SQL> -- test the function
```

```
SQL> select convert_fx('A',1000,trunc(sysdate)-3) from dual;
```

```
CONVERT_FX('A',1000,TRUNC(SYSDATE)-3)
-----
                             65617
```

```
SQL> select convert_fx('C',3000,trunc(sysdate)-8) from dual;
```

```
CONVERT_FX('C',3000,TRUNC(SYSDATE)-8)
-----
                             186576
```

```
SQL> select currency_code, conv_rate_date, conv_rate from fx_rates
```

```
    2 where (currency_code = 'A' and conv_rate_date=trunc(sysdate)-3)
```

```
    3 OR      (currency_code = 'C' and conv_rate_date=trunc(sysdate)-8)
```

```
    4 /
```

```
C CONV_RATE  CONV_RATE
- - - - -
A 26-OCT-10    65.617
C 21-OCT-10    62.192
```

```
SQL>
```

(c) Hemant K Chitale 2010

SQL in Functions

SQL>

I have created the Test Data and Function.

UPDATE using the Function:

Now, I run the Batch Update, using the Function in an UPDATE statement.

```
SQL> select /*+ FULL (f) */ count(*) from fx_rates f;
```

```
      COUNT (*)
-----
      26000
```

```
SQL> select /*+ FULL (t) */ count(*) from transactions t;
```

```
      COUNT (*)
-----
     100000
```

SQL>

SQL>

SQL> -- create a new session so that I can report session statistics

SQL> connect hemant/hemant

Connected.

SQL>

SQL> set timing on

SQL> exec dbms_session.session_trace_enable();

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.00

SQL>

SQL> update transactions

2 set converted_amount =

3

convert_fx(transaction_currency_code,transaction_amount,transaction_date)

4 /

100000 rows updated.

Elapsed: 00:00:13.77

SQL>

SQL> exec dbms_session.session_trace_disable();

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.00

SQL>

The Update of 100,000 rows took 13.77seconds.

If I run a tkprof against the trace, I get :

(c) Hemant K Chitale 2010

SQL in Functions

TKPROF: Release 11.2.0.1.0 - Development on Fri Oct 29 15:26:50 2010

Copyright (c) 1982, 2009, Oracle and/or its affiliates. All rights reserved.

Trace file: Demo_SQL_in_Function.TRC

Sort options: default

```
*****
count    = number of times OCI procedure was executed
cpu       = cpu time in seconds executing
elapsed   = elapsed time in seconds executing
disk      = number of physical reads of buffers from disk
query     = number of buffers gotten for consistent read
current   = number of buffers gotten in current mode (usually for update)
rows      = number of rows processed by the fetch or execute call
*****
```

SQL ID: abnbltj3n04cv

Plan Hash: 0

BEGIN dbms_session.session_trace_enable(); END;

call	count	cpu	elapsed	disk	query	current	rows
Parse	0	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	1
Fetch	0	0.00	0.00	0	0	0	0
total	1	0.00	0.00	0	0	0	1

Misses in library cache during parse: 0

Optimizer mode: ALL_ROWS

Parsing user id: 91

Elapsed times include waiting on following events:

Event waited on	Times Waited	Max. Wait	Total Waited
SQL*Net message to client	1	0.00	0.00
SQL*Net message from client	1	0.00	0.00

update transactions

set converted_amount =

convert_fx(transaction_currency_code,transaction_amount,transaction_date)

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	11.54	12.12	0	3284	201977	100000
Fetch	0	0.00	0.00	0	0	0	0
total	2	11.54	12.13	0	3284	201977	100000

Misses in library cache during parse: 1

Optimizer mode: ALL_ROWS

Parsing user id: 91

Rows Row Source Operation

```
-----
0  UPDATE TRANSACTIONS (cr=204326 pr=0 pw=0 time=0 us)
100000  TABLE ACCESS FULL TRANSACTIONS (cr=3275 pr=0 pw=0 time=209332 us cost=889
size=2800000 card=100000)
```

Elapsed times include waiting on following events:

Event waited on	Times Waited	Max. Wait	Total Waited
Disk file operations I/O	10	0.00	0.00
control file sequential read	42	0.00	0.00
db file sequential read	4	0.00	0.00
Data file init write	14	0.04	0.15

(c) Hemant K Chitale 2010

SQL in Functions

```
db file single write                2          0.00          0.00
control file parallel write         6          0.00          0.00
log file switch (checkpoint incomplete) 1          0.00          0.00
log buffer space                    1          0.05          0.05
SQL*Net message to client           1          0.00          0.00
SQL*Net message from client         1          0.00          0.00
*****
```

```
SQL ID: gz1jkmjq5s5fg
Plan Hash: 3275614890
SELECT :B3 *CONV_RATE
FROM
  FX_RATES WHERE CURRENCY_CODE=:B2 AND CONV_RATE_DATE=:B1
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	100000	1.18	1.09	0	0	0	0
Fetch	100000	0.57	0.54	0	201000	0	1000
total	200001	1.76	1.63	0	201000	0	1000

```
Misses in library cache during parse: 1
Misses in library cache during execute: 1
Optimizer mode: ALL_ROWS
Parsing user id: 91 (recursive depth: 1)
```

```
Rows      Row Source Operation
-----
0      TABLE ACCESS BY INDEX ROWID FX_RATES (cr=2 pr=0 pw=0 time=0 us cost=2 size=15 card=1)
0      INDEX UNIQUE SCAN MY_FX_RATES_PK (cr=2 pr=0 pw=0 time=0 us cost=1 size=0
card=1) (object id 74590)
```

```
SQL ID: 23d3sap7cask4
Plan Hash: 0
BEGIN dbms_session.session_trace_disable(); END;
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	1
Fetch	0	0.00	0.00	0	0	0	0
total	2	0.00	0.00	0	0	0	1

```
Misses in library cache during parse: 0
Optimizer mode: ALL_ROWS
Parsing user id: 91
```

OVERALL TOTALS FOR ALL NON-RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
Parse	2	0.00	0.00	0	0	0	0
Execute	3	11.54	12.12	0	3284	201977	100002
Fetch	0	0.00	0.00	0	0	0	0
total	5	11.55	12.13	0	3284	201977	100002

```
Misses in library cache during parse: 1
```

```
Elapsed times include waiting on following events:
Event waited on                      Times    Max. Wait    Total Waited
-----
SQL*Net message to client              2         0.00         0.00
SQL*Net message from client            2         0.00         0.00
```

(c) Hemant K Chitale 2010

SQL in Functions

Disk file operations I/O	10	0.00	0.00
control file sequential read	42	0.00	0.00
db file sequential read	4	0.00	0.00
Data file init write	14	0.04	0.15
db file single write	2	0.00	0.00
control file parallel write	6	0.00	0.00
log file switch (checkpoint incomplete)	1	0.00	0.00
log buffer space	1	0.05	0.05

OVERALL TOTALS FOR ALL RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
Parse	12	0.00	0.00	0	0	0	0
Execute	100011	1.18	1.09	0	0	0	0
Fetch	100020	0.57	0.54	0	201042	0	1011
total	200043	1.76	1.63	0	201042	0	1011

Misses in library cache during parse: 1
Misses in library cache during execute: 1

100003 user SQL statements in session.
11 internal SQL statements in session.
100014 SQL statements in session.

Trace file: Demo_SQL_in_Function.TRC
Trace file compatibility: 11.1.0.7
Sort options: default

1 session in tracefile.
100003 user SQL statements in trace file.
11 internal SQL statements in trace file.
100014 SQL statements in trace file.
6 unique SQL statements in trace file.
700259 lines in trace file.
13 elapsed seconds in trace file.

Wait Events : You will notice a number of Wait Events like “Disk file operations I/O”, “Data file init write”, “log buffer space”, “control file sequential read” and “control file parallel write”. These may appear occasionally in production environments depending on the pattern and nature of calls. Since the time for these Waits is not significant, they do not impact the focus of this Test Case.

Furthermore, selected Session Statistics and Waits (querying V\$SESSTAT and V\$SESSION_EVENT) are:

Statistic or Wait	Val or Cnt	Time MS
CPU : CPU used by this session	0	13,330.00
Recursive CPU : recursive cpu usage	0	8,350.00
Recursive CPU double counting: recursive cpu usage	0	-8,350.00
Statistic : consistent gets	204,374	
Statistic : execute count	100,029	
Statistic : index fetch by key	100,000	
Statistic : recursive calls	199,239	
Statistic : table fetch by rowid	1,000	
Statistic : table scan blocks gotten	3,291	
Statistic : table scan rows gotten	100,212	
Statistic : user calls	26	
Wait : Data file init write	14	159.71
Wait : Disk file operations I/O	11	.82
Wait : SQL*Net message from client	16	2.71
Wait : SQL*Net message to client	17	.02
Wait : control file parallel write	6	3.09
Wait : control file sequential read	42	.27
Wait : db file sequential read	4	.03
Wait : db file single write	2	.44

SQL in Functions

Wait : events in waitclass Other	1	.00
Wait : log buffer space	1	55.64
Wait : log file switch (checkpoint incomplete)	1	6.30
Wait : log file sync	1	.24

sum		13,559.26

These are the critical things to note:

1. OF the total elapsed time of 13.77seconds, CPU Usage was 13.33seconds (tkprof reports a total CPU time of 13.31seconds) of which 8.35seconds were on recursive calls – the repeated calls to the Function, invoking an SQL for each row. Therefore, a very large portion of the total elapsed time was because Oracle had to use CPU time for the recursive calls. (The difference between 13.77seconds and 13.56 in the Statistics + Waits is because of rounding in computation of CPU and Waits. This difference is not significant).
2. Although the “update transactions” was executed only once, a “SELECT :B3 *CONV_RATE FROM FX_RATES WHERE CURRENCY_CODE=:B2 AND CONV_RATE_DATE=:B1” was executed 100,000 times. (Note how Oracle rewrites SQL in a PLSQL block (the CONVERT_FX function in this case) so that it is all uppercase and with Binds – this is how Oracle makes the SQL reusable and not requiring a hard parse at each execute).
3. The “recursive calls” statistic of 199,239 and “execute count” statistic of 100,029 are another indication of the number of context switches Oracle had to make between PLSQL and SQL. (I have excluded SYS calls from the tkprof listing).

Using a Correlated UPDATE:

Here is an alternative way of running the UPDATE, as a Correlated Update in SQL *without invoking the Function itself*:

```
SQL>
SQL> select /*+ FULL (f) */ count(*) from fx_rates f;

COUNT (*)
-----
      26000

SQL> select /*+ FULL (t) */ count(*) from transactions t;

COUNT (*)
-----
     100000

SQL>
SQL> connect hemant/hemant
Connected.
SQL> exec dbms_session.session_trace_enable();
```

SQL in Functions

PL/SQL procedure successfully completed.

SQL>

SQL> set timing on

SQL> REM This is the SQL Update -- use a Correlated Subquery on the PrimaryKey

```
SQL> update transactions mt
  2  set mt.converted_amount=
  3    (select st.transaction_amount*sf.conv_rate
  4      from transactions st, fx_rates sf
  5      where
  6        st.transaction_currency_code=sf.currency_code
  7      and
  8        st.transaction_date=sf.conv_rate_date
  9      and st.transaction_id=mt.transaction_id),
 10  mt.updated_at=sysdate;
```

100000 rows updated.

Elapsed: 00:00:04.22

SQL>

SQL> commit;

Commit complete.

Elapsed: 00:00:00.00

SQL> exec dbms_session.session_trace_disable();

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.01

SQL>

This Update took 4.22seconds.

The tkprof of the SQL Correlated Update shows:

TKPROF: Release 11.2.0.1.0 - Development on Fri Oct 29 15:31:42 2010

Copyright (c) 1982, 2009, Oracle and/or its affiliates. All rights reserved.

Trace file: Run_Correlated_Update.TRC

Sort options: default

```
*****
count      = number of times OCI procedure was executed
cpu        = cpu time in seconds executing
elapsed    = elapsed time in seconds executing
disk       = number of physical reads of buffers from disk
query      = number of buffers gotten for consistent read
current    = number of buffers gotten in current mode (usually for update)
rows       = number of rows processed by the fetch or execute call
*****
```

SQL ID: abnbltj3n04cv

Plan Hash: 0

BEGIN dbms_session.session_trace_enable(); END;

call	count	cpu	elapsed	disk	query	current	rows
Parse	0	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	1

(c) Hemant K Chitale 2010

SQL in Functions

Fetch	0	0.00	0.00	0	0	0	0
total	1	0.00	0.00	0	0	0	1

Misses in library cache during parse: 0
Optimizer mode: ALL_ROWS
Parsing user id: 91

Elapsed times include waiting on following events:

Event waited on	Times Waited	Max. Wait	Total Waited
SQL*Net message to client	1	0.00	0.00
SQL*Net message from client	1	0.00	0.00

```
update transactions mt
set mt.converted_amount=
(select st.transaction_amount*sf.conv_rate
 from transactions st, fx_rates sf
 where
  st.transaction_currency_code=sf.currency_code
 and
  st.transaction_date=sf.conv_rate_date
 and st.transaction_id=mt.transaction_id),
mt.updated_at=sysdate
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	3.65	4.21	0	429159	212502	100000
Fetch	0	0.00	0.00	0	0	0	0
total	2	3.65	4.21	0	429159	212502	100000

Misses in library cache during parse: 1
Optimizer mode: ALL_ROWS
Parsing user id: 91

Rows Row Source Operation

```
-----
0 UPDATE TRANSACTIONS (cr=417584 pr=0 pw=0 time=0 us)
202868 TABLE ACCESS FULL TRANSACTIONS (cr=6645 pr=0 pw=0 time=308451 us cost=889
size=1900000 card=100000)
1027 NESTED LOOPS (cr=422490 pr=0 pw=0 time=0 us cost=3 size=35 card=1)
102868 TABLE ACCESS BY INDEX ROWID TRANSACTIONS (cr=215727 pr=0 pw=0 time=0 us cost=2
size=20 card=1)
102868 INDEX UNIQUE SCAN MY_TRANSACTIONS_PK (cr=112859 pr=0 pw=0 time=0 us cost=1 size=0
card=1)(object id 74592)
1027 TABLE ACCESS BY INDEX ROWID FX_RATES (cr=206763 pr=0 pw=0 time=0 us cost=1
size=390000 card=26000)
1027 INDEX UNIQUE SCAN MY_FX_RATES_PK (cr=205736 pr=0 pw=0 time=0 us cost=0 size=0
card=1)(object id 74590)
```

Elapsed times include waiting on following events:

Event waited on	Times Waited	Max. Wait	Total Waited
Disk file operations I/O	16	0.00	0.00
reliable message	2	0.00	0.00
rdbms ipc reply	1	0.00	0.00
control file sequential read	84	0.01	0.01
db file sequential read	8	0.00	0.00
Data file init write	28	0.07	0.24
db file single write	4	0.00	0.00
control file parallel write	12	0.00	0.00
log file switch completion	2	0.00	0.00
log buffer space	1	0.30	0.30
SQL*Net message to client	1	0.00	0.00
SQL*Net message from client	1	0.00	0.00

SQL ID: 23wm3kz7rps5y
Plan Hash: 0
commit

(c) Hemant K Chitale 2010

SQL in Functions

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	1	0
Fetch	0	0.00	0.00	0	0	0	0
total	2	0.00	0.00	0	0	1	0

Misses in library cache during parse: 0
Parsing user id: 91

Elapsed times include waiting on following events:

Event waited on	Times Waited	Max. Wait	Total Waited
SQL*Net message to client	1	0.00	0.00
SQL*Net message from client	1	0.00	0.00

SQL ID: 23d3sap7cask4
Plan Hash: 0
BEGIN dbms_session.session_trace_disable(); END;

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	1
Fetch	0	0.00	0.00	0	0	0	0
total	2	0.00	0.00	0	0	0	1

Misses in library cache during parse: 0
Optimizer mode: ALL_ROWS
Parsing user id: 91

OVERALL TOTALS FOR ALL NON-RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
Parse	3	0.00	0.00	0	0	0	0
Execute	4	3.65	4.21	0	429159	212503	100002
Fetch	0	0.00	0.00	0	0	0	0
total	7	3.66	4.21	0	429159	212503	100002

Misses in library cache during parse: 1

Elapsed times include waiting on following events:

Event waited on	Times Waited	Max. Wait	Total Waited
SQL*Net message to client	3	0.00	0.00
SQL*Net message from client	3	0.00	0.00
Disk file operations I/O	16	0.00	0.00
reliable message	2	0.00	0.00
rdbms ipc reply	1	0.00	0.00
control file sequential read	84	0.01	0.01
db file sequential read	8	0.00	0.00
Data file init write	28	0.07	0.24
db file single write	4	0.00	0.00
control file parallel write	12	0.00	0.00
log file switch completion	2	0.00	0.00
log buffer space	1	0.30	0.30

OVERALL TOTALS FOR ALL RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
------	-------	-----	---------	------	-------	---------	------

SQL in Functions

Parse	28	0.00	0.00	0	0	0	0
Execute	28	0.00	0.00	0	0	0	0
Fetch	51	0.00	0.00	0	107	0	28
-----	-----	-----	-----	-----	-----	-----	-----
total	107	0.00	0.00	0	107	0	28

Misses in library cache during parse: 0

```

    4 user SQL statements in session.
    28 internal SQL statements in session.
    32 SQL statements in session.
*****
Trace file: Run_Correlated_Update.TRC
Trace file compatibility: 11.1.0.7
Sort options: default

```

```

    1 session in tracefile.
    4 user SQL statements in trace file.
    28 internal SQL statements in trace file.
    32 SQL statements in trace file.
    6 unique SQL statements in trace file.
   486 lines in trace file.
    4 elapsed seconds in trace file.

```

Selected Session Statistics and Waits for the Correlated Update are:

Statistic or Wait	Val or Cnt	Time MS
CPU : CPU used by this session	0	3,670.00
Recursive CPU : recursive cpu usage	0	100.00
Recursive CPU double counting: recursive cpu usage	0	-100.00
Statistic : consistent gets	429,314	
Statistic : execute count	46	
Statistic : index fetch by key	205,736	
Statistic : recursive calls	387	
Statistic : table fetch by rowid	103,895	
Statistic : table scan blocks gotten	6,681	
Statistic : table scan rows gotten	203,304	
Statistic : user calls	28	
Wait : Data file init write	28	245.19
Wait : Disk file operations I/O	17	1.31
Wait : SQL*Net message from client	17	3.02
Wait : SQL*Net message to client	18	.02
Wait : control file parallel write	12	7.09
Wait : control file sequential read	84	13.54
Wait : db file sequential read	8	.07
Wait : db file single write	4	.95
Wait : events in waitclass Other	4	.08
Wait : log buffer space	1	305.12
Wait : log file switch completion	2	9.78
-----	-----	-----
sum		4,256.16

The Correlated UPDATE ran in 4.22seconds --- much faster than the using the Function, which took 13.77seconds.

These are the critical things to note:

1. CPU time was only 3.66 or 3.67 seconds.
2. However, “consistent gets” is slightly more than twice that in the first method – at being in excess of 420thousand. Most of these are because of the index lookups. (Similarly “index fetch by key” is twice that in the first method). The TRANSACTIONS table is visited twice and indexed lookups on FX_RATES are in a

Nested Loop. Can the Correlated Update be improved further? You could take that as a further exercise.

3. The “recursive calls” is only 387 and “execute count” is only 46. Much lower than the 199thousand and 100thousand calls of the first method.

A Comparison of the Two Tests:

Statistic\TestCase	Update using Function	SQL Correlated Update	Remarks
Elapsed Time	13.77seconds	4.22seconds	
CPU Time	13.33seconds	03.67seconds	
Recursive CPU Time	08.35seconds	00.10seconds	This is where the major portion of the difference in elapsed time occurs
“recursive calls”	199,239	387	These cause additional for CPU usage
“execute count”	100,029	046	
“consistent gets”	204,374	429,314	This is where the Correlated Update needs improvement
“index fetch by key”	100,000	205,736	
“table fetch by rowid”	1,000	103,895	

Key Findings in this Test Case:

1. Context Switches caused by repeated SQL calls in PLSQL increase the CPU Utilisation. Significantly!
2. Function Calls that run SELECT statements (that retrieve more than a few blocks) can add overhead when executed for *each row*.

Recommendations:

1. Recognise the difference between PLSQL "Procedural" operations and SQL "Set" Operations. Developers who come from backgrounds with other programming languages (eg Java) typically overuse Procedural operations. It is more important to learn SQL first and well.
2. Functions are fine when calling sparingly for a few rows – e.g. in data validation or input screens in an OLTP environment or when operating on small data sets. If, however, a

Function is being invoked repeatedly for each row in the ResultSet, review the impact of the recursive calls to the SQL(s) in the Function.

3. Example of Functions that do not do lookups on tables are the inbuilt 'SYSDATE', 'TO_CHAR' and other functions, which themselves do not run SELECTs on table blocks.

4. Use SQL-only code wherever possible.

When SQL can execute an operation on multiple rows in one single call, it is best to use SQL. (See Tom Kyte's "*mantra*"

http://asktom.oracle.com/pls/asktom/f?p=100:11:0:::::P11_QUESTION_ID:760210800346068768)